



Files and Directories

Hongik University

Eunsung Jung

Disclaimer: The slides are borrowed from many sources!

stat(2) family of functions

```
#include <sys/types.h>
#include <sys/stat.h>

int stat(const char *path, struct stat *sb);
int lstat(const char *path, struct stat *sb);
int fstat(int fd, struct stat *sb);
```

Returns: 0 if OK, -1 on error

- All these functions return extended attributes about the referenced file (in the case of symbolic links, lstat(2) returns attributes of the link, others return stats of the referenced file).
- Get attribute info using **struct stat**



Data structure: struct stat

What is this?

Try `ls -l /dev/sda``

```
struct stat {
    dev_t      st_dev;      /* device number (filesystem) */
    ino_t      st_ino;      /* i-node number (serial number) */
    mode_t     st_mode;     /* file type & mode (permissions) */
    dev_t      st_rdev;     /* device number for special files */
    nlink_t    st_nlink;    /* number of links */
    uid_t      st_uid;      /* user ID of owner */
    gid_t      st_gid;      /* group ID of owner */
    off_t      st_size;     /* size in bytes, for regular files */
    time_t     st_atime;    /* time of last access */
    time_t     st_mtime;    /* time of last modification */
    time_t     st_ctime;    /* time of last file status change */
    long       st_blocks;   /* number of 512-byte* blocks allocated */
    long       st_blksize;  /* best I/O block size */
};
```



How to identify file types?

- The **st_mode** field of the **struct stat** encodes the type of file:
 - **regular** – most common, interpretation of data is up to application
 - **directory** – contains names of other files and pointer to information on those files. Any process can read, only kernel can write.
 - **character special** – used for character types of devices
 - **block special** – used for disk devices (typically). All devices are either character or block special.
 - **FIFO** – used for interprocess communication (sometimes called named pipe)
 - **socket** – used for network communication and non-network communication (same host).
 - **symbolic link** – Points to another file.
- Find out more in `<sys/stat.h>`, `<asm/stat.h>`.



How to identify file types?

Macro	Type of file
<code>S_ISREG()</code>	regular file
<code>S_ISDIR()</code>	directory file
<code>S_ISCHR()</code>	character special file
<code>S_ISBLK()</code>	block special file
<code>S_ISFIFO()</code>	pipe or FIFO
<code>S_ISLNK()</code>	symbolic link
<code>S_ISSOCK()</code>	socket

Figure 4.1 File type macros in `<sys/stat.h>`



Figure 4.3: Print type of file

```
int
main(int argc, char *argv[])
{
    int          i;
    struct stat buf;
    char         *ptr;

    for (i = 1; i < argc; i++) {
        printf("%s: ", argv[i]);
        if (lstat(argv[i], &buf) < 0) {
            err_ret("lstat error");
            continue;
        }
        if (S_ISREG(buf.st_mode))
            ptr = "regular";
        else if (S_ISDIR(buf.st_mode))
            ptr = "directory";
        else if (S_ISCHR(buf.st_mode))
            ptr = "character special";
        else if (S_ISBLK(buf.st_mode))
            ptr = "block special";
        else if (S_ISFIFO(buf.st_mode))
            ptr = "fifo";
        else if (S_ISLNK(buf.st_mode))
            ptr = "symbolic link";
        else if (S_ISSOCK(buf.st_mode))
            ptr = "socket";
        else
            ptr = "** unknown mode **";
        printf("%s\n", ptr);
    }
    exit(0);
}
```



Figure 4.3: Results

```
$ ./a.out /etc/passwd /etc /dev/log /dev/tty \  
> /var/lib/oprofile/opd_pipe /dev/sr0 /dev/cdrom  
/etc/passwd: regular  
/etc: directory  
/dev/log: socket  
/dev/tty: character special  
/var/lib/oprofile/opd_pipe: fifo  
/dev/sr0: block special  
/dev/cdrom: symbolic link
```

- We can make a FIFO(named pipe)
 - mkfifo /tmp/testpipe



Set-User-ID and Set-Group-ID

- **Every process** has six or more IDs associated with it:

real user ID real group ID	who we really are
effective user ID effective group ID supplementary group IDs	used for file access permission checks
saved set-user-ID saved set-group-ID	saved by <code>exec</code> functions

Figure 4.5 User IDs and group IDs associated with each process

- **Real user/group ID:** Normally, these values **don't change** during a login session.
- **Effective user/group ID and supplementary group IDs:** determine our file access permissions. **Normally, the effective user ID equals the real user ID, and the effective group ID equals the real group ID.**
- **Saved set-user-ID and saved set-group-ID:** contain copies of the effective user ID and the effective group ID, respectively, when a program is executed.



set-user-ID bit/set-group-ID bit

- Process user ID/group ID vs. File user ID/group ID
 - The real user ID can be obtained by calling `getuid()`.
 - If set-user-ID bit and set-group-ID bit are set, **effective user ID would be root.**



File Access Permissions

- Nine permission bits for each file

st_mode mask	Meaning
S_IRUSR	user-read
S_IWUSR	user-write
S_IXUSR	user-execute
S_IRGRP	group-read
S_IWGRP	group-write
S_IXGRP	group-execute
S_IROTH	other-read
S_IWOTH	other-write
S_IXOTH	other-execute

Figure 4.6 The nine file access permission bits, from `<sys/stat.h>`



File Access Permissions

- r vs. x permission of a directory
 - x bit for a directory is the search bit.
 - For example, to open the file `/usr/include/stdio.h` , we need execute permission in the directory `/` , execute permission in the directory `/usr` , and execute permission in the directory `/usr/include` .
 - r bit for a directory is the read bit for the files under that directory.
 - Read permission lets us read the directory, obtaining a list of all the filenames in the directory.



File Access Permissions

■ Summary

- To open a file, need execute permission on each directory component of the path.
- To open a file with `O_RDONLY` or `O_RDWR`, need read permission.
- To open a file with `O_WRONLY` or `O_RDWR`, need write permission.
- To use `O_TRUNC`, must have write permission.
- To create a new file, must have write+execute permission for the directory.
- To delete a file, need write+execute on directory, file doesn't matter
- To execute a file (via `exec` family), need execute permission



File Access Permission Test by Kernel

1. If `effective-uid == 0`, grant access
2. If `effective-uid == st_uid`
 - 1) if appropriate user permission bit is set, grant access
 - 2) else, deny access
3. If `effective-gid == st_gid`
 - 1) if appropriate group permission bit is set, grant access
 - 2) else, deny access
4. If appropriate other permission bit is set, grant access, else deny access



Ownership of new files and directories

- `st_uid` = effective-uid
- `st_gid` = ...either:
 - effective-gid of process.
 - gid of directory in which it is being created.



access(2)

```
#include <unistd.h>
```

```
int access(const char *pathname, int mode);
```

```
int faccessat(int fd, const char *pathname, int mode, int flag);
```

Both return: 0 if OK, -1 on error

- Tests file accessibility **on the basis of the real uid and gid**. Allows setuid/setgid programs to see if the real user could access the file without it having to drop permissions to do so.

<i>mode</i>	Description
R_OK	test for read permission
W_OK	test for write permission
X_OK	test for execute permission

Figure 4.7 The *mode* flags for access function, from <unistd.h>



umask(2)

```
#include <sys/stat.h>
```

```
mode_t umask(mode_t numask);
```

Returns: previous file mode creation mask

- umask(2) sets the **file creation mode mask**. Any bits that are on in the file creation mask are turned off in the file's mode.
- If a program needs to be able to ensure certain permissions on a file, it may need to turn off (or modify) the umask, **which affects only the current process**.



Example

```
#include "apue.h"
#include <fcntl.h>

#define RWRWRW (S_IRUSR|S_IWUSR|S_IRGRP|S_IWGRP|S_IROTH|S_IWOTH)

int
main(void)
{
    umask(0);
    if (creat("foo", RWRWRW) < 0)
        err_sys("creat error for foo");
    umask(S_IRGRP | S_IWGRP | S_IROTH | S_IWOTH);
    if (creat("bar", RWRWRW) < 0)
        err_sys("creat error for bar");
    exit(0);
}
```

disables all the group and other permission bits.

Figure 4.9 Example of umask function

```
$ umask
002
$ ./a.out
$ ls -l foo bar
-rw----- 1 sar
-rw-rw-rw- 1 sar
$ umask
002
```

first print the current file mode creation mask

```
0 Dec  7 21:20 bar
0 Dec  7 21:20 foo
```

see if the file mode creation mask changed



chmod(2), lchmod(2) and fchmod(2):

```
#include <sys/stat.h>
```

```
int chmod(const char *pathname, mode_t mode);
```

```
int fchmod(int fd, mode_t mode);
```

```
int fchmodat(int fd, const char *pathname, mode_t mode, int flag);
```

All three return: 0 if OK, -1 on error

- Changes the permissions of an existing file.



chmod(2), lchmod(2) and fchmod(2):

<i>mode</i>	Description
S_ISUID	set-user-ID on execution
S_ISGID	set-group-ID on execution
S_ISVTX	saved-text (sticky bit)
S_IRWXU	read, write, and execute by user (owner)
S_IRUSR	read by user (owner)
S_IWUSR	write by user (owner)
S_IXUSR	execute by user (owner)
S_IRWXG	read, write, and execute by group
S_IRGRP	read by group
S_IWGRP	write by group
S_IXGRP	execute by group
S_IRWXO	read, write, and execute by other (world)
S_IROTH	read by other (world)
S_IWOTH	write by other (world)
S_IXOTH	execute by other (world)

Figure 4.11 The *mode* constants for chmod functions, from `<sys/stat.h>`

■ Sticky Bit

- If it was set for an executable program file, then the first time the program was executed, a copy of the program's text was saved in the **swap** area when the process terminated.
- Next time, the program will be loaded into memory more quickly.
- **Not supported by Linux.**



chown(2), lchown(2) and fchown(2)

```
#include <unistd.h>
```

```
int chown(const char *path, uid_t owner, gid_t group);  
int lchown(const char *path, uid_t owner, gid_t group);  
int fchown(int fd, uid_t owner, gid_t group);
```

Returns: 0 if OK, -1 on error

- Changes `st_uid` and `st_gid` for a file. For BSD, must be superuser.
- owner or group can be -1 to indicate that it should remain the same.
- Non-superusers can change the `st_gid` field if both:
 - effective-user ID == `st_uid` and
 - owner == file's user ID and group == effective-group ID (or one of the supplementary group IDs)



File Size

- `st_size` member of the struct `stat`.
- For a regular file, a file size of 0 is allowed.
- For a symbolic link, the file size is
 - The number of bytes in the file name.
 - E.g. `lrwxrwxrwx 1 root 7 Sep 25 07:14 lib -> usr/lib`
- Holes in a File

```
$ ls -l core
-rw-r--r--  1 sar 8483248 Nov 18 12:18 core
$ du -s core
272      core
```



truncate(2) and ftruncate(2)

```
#include <unistd.h>
```

```
int truncate(const char *pathname, off_t length);
```

```
int ftruncate(int fd, off_t length);
```

Both return: 0 if OK, -1 on error

- Truncate an existing file to *length* bytes.

