



Files and Directories

Hongik University

Eunsung Jung

Disclaimer: The slides are borrowed from many sources!

File System Overview

- A disk can be divided into logical *partitions*.
- Each partition can contain a *file system*.
- A file system is divided into *multiple cylinder groups in general*.
 - *Different file systems have different structures.*

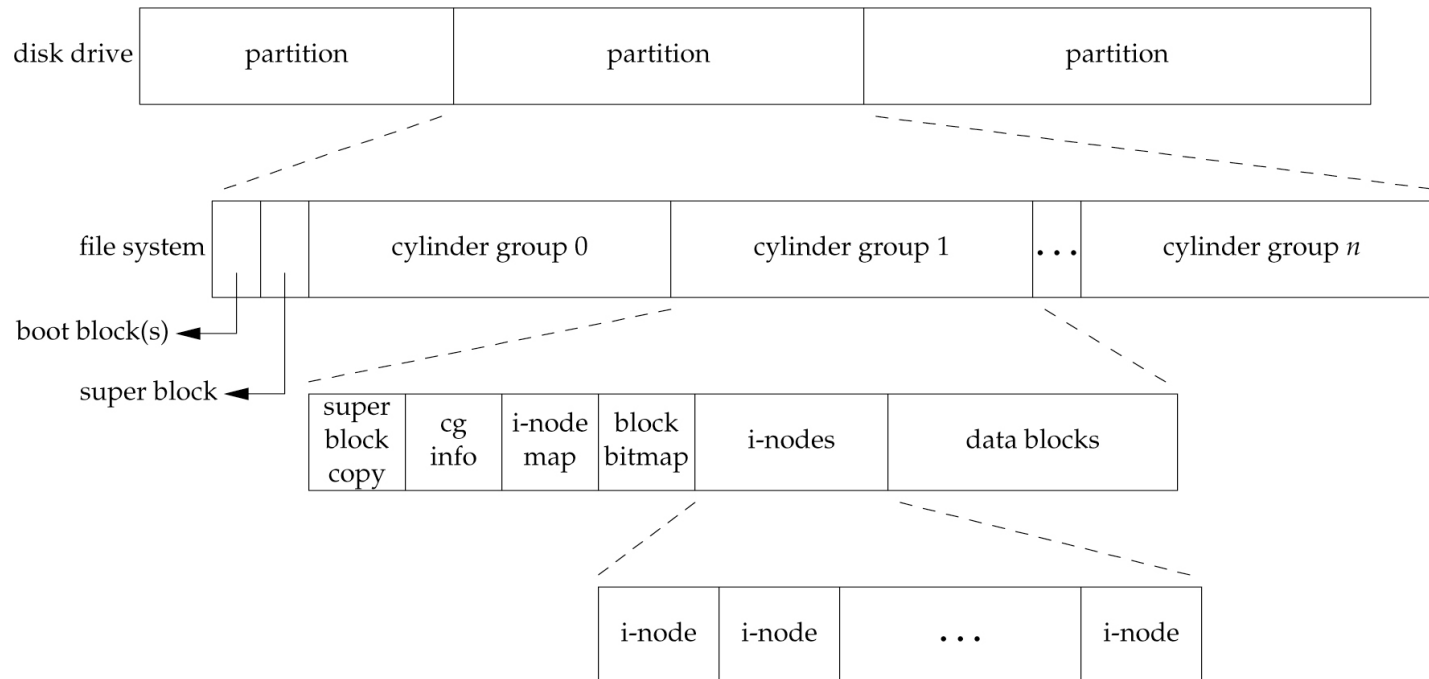


Figure 4.13 Disk drive, partitions, and a file system

Cylinder Group

- Each cylinder group contains *a list of inodes (i-list)* as well as the actual directory- and data blocks
- A directory entry is really just a hard link *mapping a “filename” to an inode*.

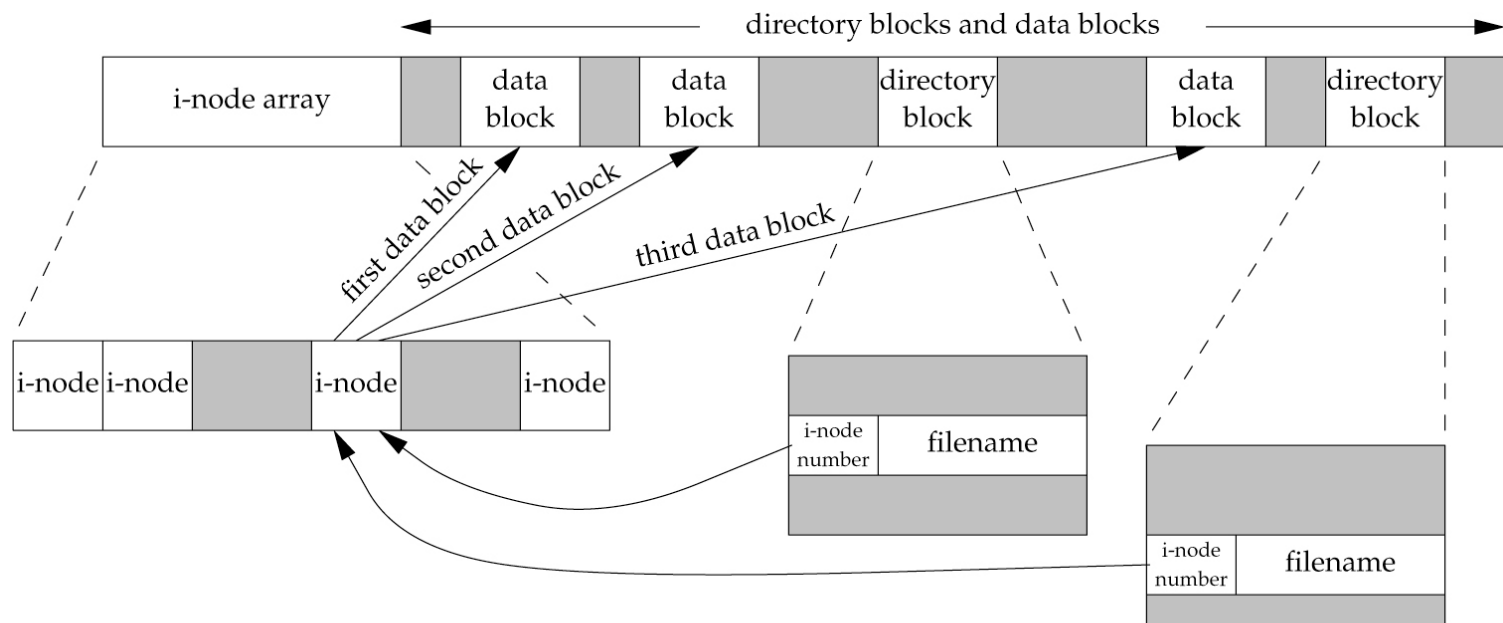


Figure 4.14 Cylinder group's i-nodes and data blocks in more detail



Hard link vs Symbolic link

- Hard link: Multiple directory entries can point to the same i-node. This is tracked by *inode count*, and the i-node can be delete only if the i-node count is 0.
- Symbolic link: The contents of the file stores the path.

```
lrwxrwxrwx  1 root      7 Sep 25 07:14 lib -> usr/lib
```

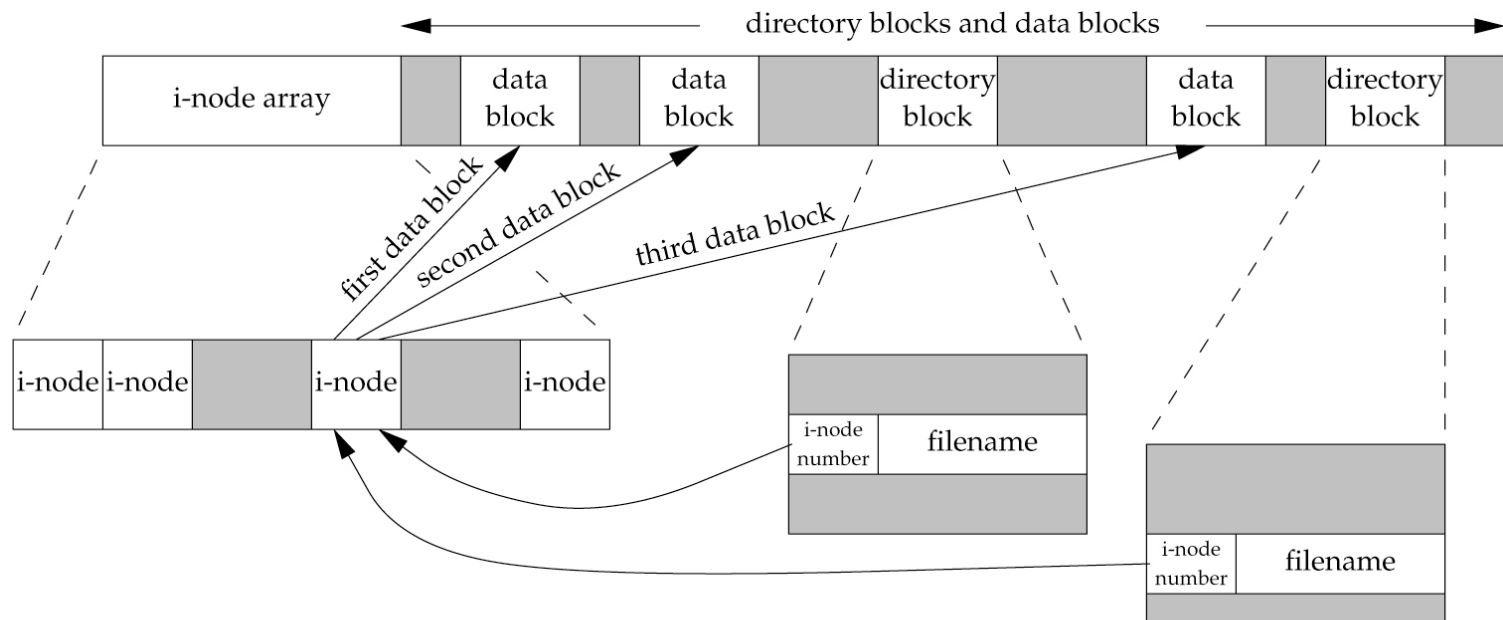


Figure 4.14 Cylinder group's i-nodes and data blocks in more detail

Directories

- Each directory contains at least two entries:
 - . (this directory), .. (the parent directory)
- the link count (st_nlink) of a directory is at least 2: **i-node 2549** has . and *directory name* link.

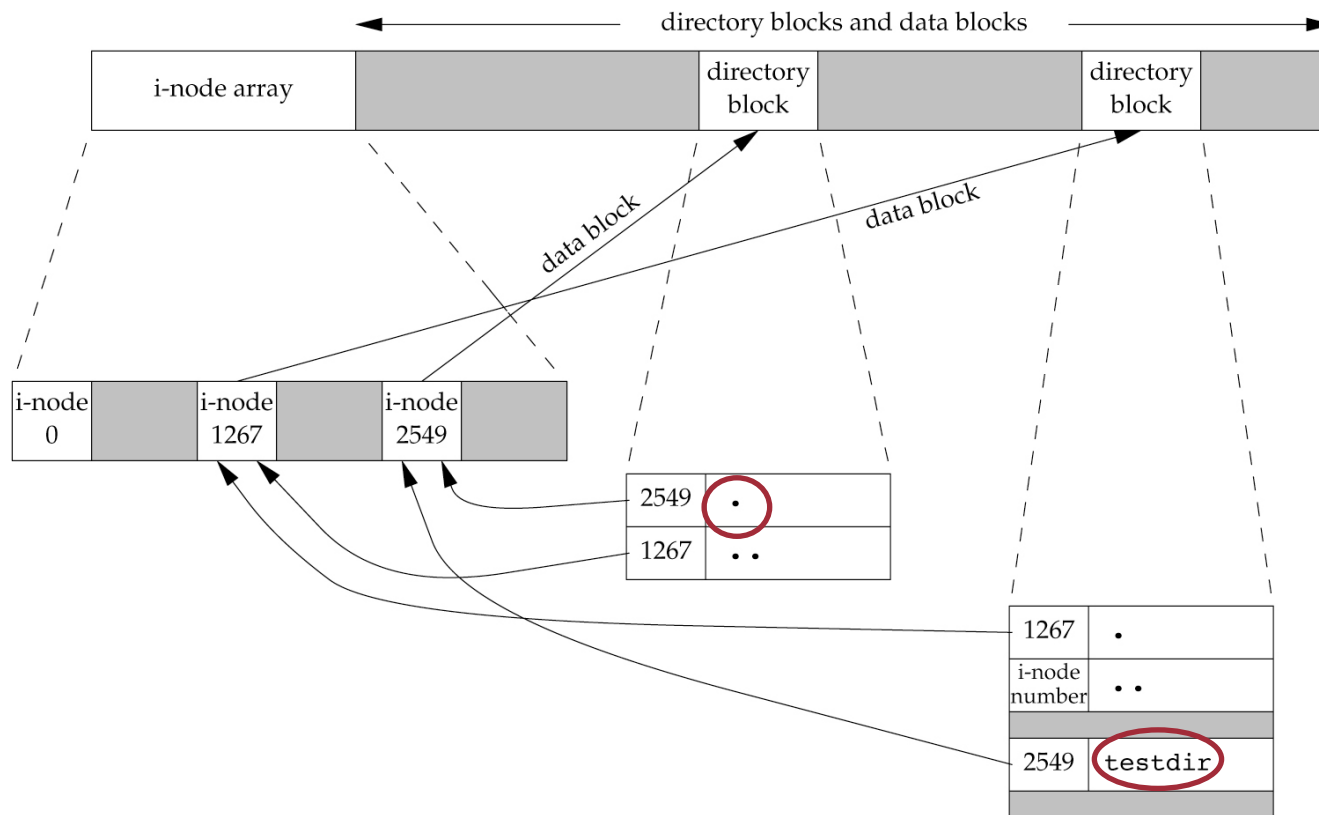


Figure 4.15 Sample cylinder group after creating the directory `testdir`

Inodes

- The **inode** contains most of the information found in the stat structure.
- Every inode has a link count (st_nlink): it shows how many “things” point to this inode. **Only if this link count is 0** (and no process has the file open) are the data blocks freed.
- Inode number in a directory entry must point to an inode on the same file system (**no hardlinks across filesystems**)
- To move a file within a single filesystem, we can **just “move” the directory entry (actually done by creating a new entry, and deleting the old one).**



link(2) and linkat(2)

```
#include <unistd.h>
```

```
int link(const char *existingpath, const char *newpath);
```

```
int linkat(int efd, const char *existingpath, int nfd, const char *newpath,  
           int flag);
```

Both return: 0 if OK, -1 on error

- Creates a link to an existing file (hard link).
- POSIX.1 allows links to cross filesystems, most implementations (SVR4, BSD) don't.
- Only **superuser** can create links to directories (loops in filesystem are bad)



unlink(2) and unlinkat(2)

```
#include <unistd.h>
```

```
int unlink(const char *pathname);
```

```
int unlinkat(int fd, const char *pathname, int flag);
```

Both return: 0 if OK, -1 on error

- Removes directory entry and decrements link count of file
- If file link count == 0, free data blocks associated with file (unless processes have the file open)



rename(2) and renameat(2)

```
#include <stdio.h>

int rename(const char *oldname, const char *newname);

int renameat(int oldfd, const char *oldname, int newfd,
             const char *newname);
```

Both return: 0 if OK, -1 on error

- If *oldname* refers to a file:
 - if *newname* exists and it is not a directory, it's removed and *oldname* is renamed *newname*.
 - if *newname* exists and it is a directory, an error results.
 - must have w+x perms for the directories containing *old/newname*.
- If *oldname* refers to a directory:
 - if *newname* exists and is an empty directory (contains only . and ..), it is removed; *oldname* is renamed *newname*.
 - if *newname* exists and is a file, an error results.
 - if *oldname* is a prefix of *newname* an error results.
 - must have w+x perms for the directories containing *old/newname*.



Symbolic Links

```
#include <unistd.h>
```

```
int symlink(const char *name1, const char *name2);
```

Returns: 0 if OK, -1 on error

- File whose “data” is a path to another file.
- Anyone can create symlinks to directories or files



Symbolic Links

- Certain functions **dereference** the link, others **operate** on the link.

Function	Does not follow symbolic link	Follows symbolic link
access		•
chdir		•
chmod		•
chown		•
creat		•
exec		•
lchown	•	
link		•
lstat	•	
open		•
opendir		•
pathconf		•
readlink	•	
remove	•	
rename	•	
stat		•
truncate		•
unlink	•	

Figure 4.17 Treatment of symbolic links by various functions



Symbolic Links

- How do we get the contents of a symlink? `open(2)` and `read(2)`?
 - `readlink()`

```
#include <unistd.h>
```


```
int readlink(const char *path, char *buf, size_t bufsize);
```

Returns: number of bytes placed into buffer if OK, -1 on error



Symbolic Links: Loop

- A simple program following the link will print the following.
 - To break the loop, use [unlink](#).



```
foo
foo/a
foo/testdir
foo/testdir/a
foo/testdir/testdir
foo/testdir/testdir/a
foo/testdir/testdir/testdir
foo/testdir/testdir/testdir/a
(many more lines until we encounter an ELOOP error)
```

Figure 4.18 Symbolic link `testdir` that creates a loop



File Times

- Three time values associated with each file

Field	Description	Example	ls(1) option
st_atim	last-access time of file data	read	-u
st_mtim	last-modification time of file data	write	default
st_ctim	last-change time of i-node status	chmod, chown	-c

Figure 4.19 The three time values associated with each file



File Times: Change File Times

```
#include <sys/types.h>
```

```
int utimes(const char *path, const struct timeval times[2]);  
int lutimes(const char *path, const struct timeval times[2]);  
int futimes(int fd, const struct timeval times[2]);
```

Returns: 0 if OK, -1 on error

- If times is **NULL**, **access time and modification time** are set to the current time (must be owner of file or have write permission).
- If times is **non-NULL**, then times are set according to the timeval struct array. For this, you **must be the owner of the file** (write permission not enough).



mkdir(2)

```
#include <sys/types.h>
#include <sys/stat.h>

int mkdir(const char *path, mode_t mode);
```

Returns: 0 if OK, -1 on error

- Creates a new, empty (except for . and .. entries) directory. Access permissions specified by mode and restricted by the `umask(2)` of the calling process.
- Solaris 10 and Linux 3.2.0 also have the new directory inherit the [set-group-ID bit](#) from the parent directory. Files created in the new directory will then [inherit the group ID](#) of that directory. With Linux, the file system implementation determines whether this behavior is supported.



rmkdir(2)

```
#include <unistd.h>

int mkdir(const char *path);
```

Returns: 0 if OK, -1 on error

- If the link count is 0 (after this call), and no other process has the directory open, directory is removed. Directory must be empty (only . and .. remaining)



Reading Directories

```
#include <sys/types.h>
#include <dirent.h>

DIR *opendir(const char *filename);
                                Returns: pointer if OK, NULL on error

struct dirent *readdir(DIR *dp);
                                Returns: pointer if OK, NULL at end of dir or on error

void rewinddir(DIR *dp);
int closedir(DIR *dp);
                                Returns: 0 if OK, -1 on error
```

- Read by anyone with read permission on the directory
- Format of directory is implementation dependent (always use readdir and friends)
- rewinddir resets an open directory to the beginning so readdir will again return the first entry.



Moving around directories

```
#include <unistd.h>
```

```
char *getcwd(char *buf, size_t size);
```

Returns: *buf* if OK, NULL on error

- Get the process's current working directory.

```
#include <unistd.h>
```

```
int chdir(const char *path);
```

```
int fchdir(int fd);
```

Returns: 0 if OK, -1 on error

- Allows a process to change its current working directory. **Note that `chdir` and `fchdir` affect only the current process.**



Device Special Files

- Every file system is known by its **major** and **minor** device numbers, which are encoded in the primitive system data type `dev_t`.
 - The major number identifies the device driver.
 - The minor number identifies the specific subdevice.

```
$ ./a.out / /home/sar /dev/tty[01]
/: dev = 8/3
/home/sar: dev = 8/4
/dev/tty0: dev = 0/5 (character) rdev = 4/0
/dev/tty1: dev = 0/5 (character) rdev = 4/1
$ mount which directories are mounted on which devices?
/dev/sda3 on / type ext3 (rw,errors=remount-ro,commit=0)
/dev/sda4 on /home type ext3 (rw,commit=0)
$ ls -l /dev/tty[01] /dev/sda[34]
brw-rw---- 1 root      8, 3 2011-07-01 11:08 /dev/sda3
brw-rw---- 1 root      8, 4 2011-07-01 11:08 /dev/sda4
crw--w---- 1 root      4, 0 2011-07-01 11:08 /dev/tty0
crw----- 1 root      4, 1 2011-07-01 11:08 /dev/tty1
```



Lab#4

- Figure 4.22 Recursively descend a directory hierarchy, counting file types.

