



# Process Programming (Chap7-Chap9)

Hongik University

Eunsung Jung

*Disclaimer: The slides are borrowed from many sources!*

# Overview

1. What is a Process?
2. `fork()`
3. `exec()`
4. `wait()`
5. File Descriptors across Processes
6. Special Exit Cases
7. IO Redirection



# What makes up a Process?

- program code
- machine registers
- global data
- stack
- open files (file descriptors)
- an environment (environment variables; credentials for security)



# Some of the Context Information

- Process ID (`pid`)                      unique integer
- Parent process ID (`ppid`)
- Real User ID                              ID of user/process which started this process
- Effective User ID                        ID of user who wrote the process' program
- Current directory
- File descriptor table
- Environment                              VAR=VALUE pairs



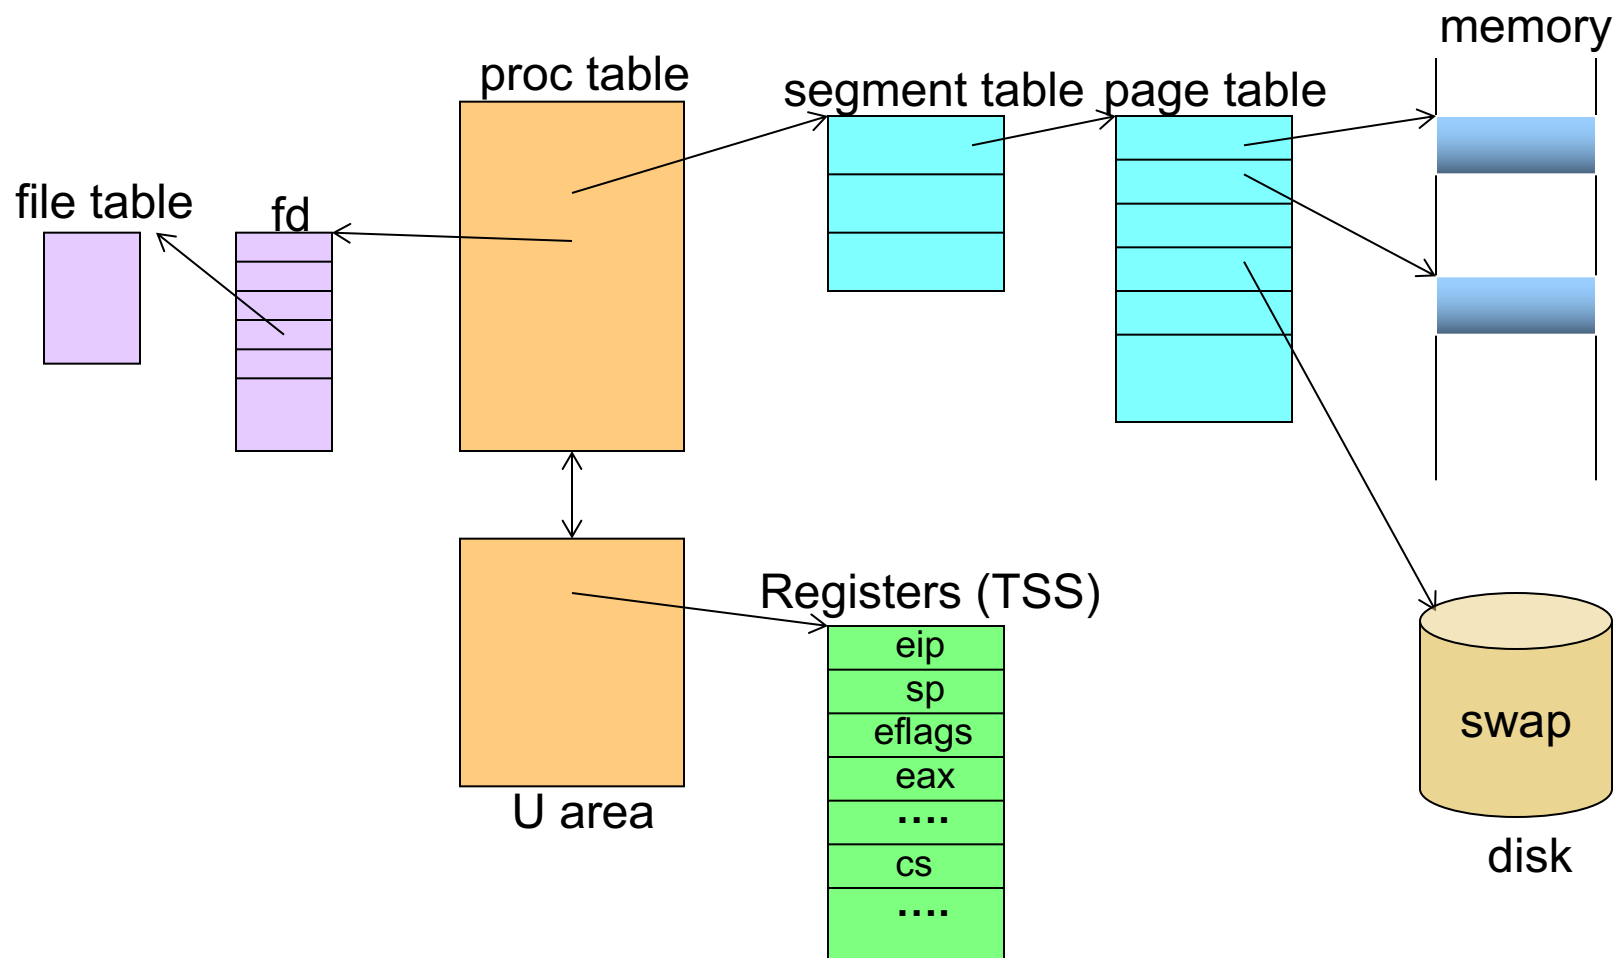
- Pointer to program code
  - Pointer to data
  - Pointer to stack
  - Pointer to heap
- 
- Execution priority
  - Signal information

Memory for global vars  
Memory for local vars  
Dynamically allocated



# Context

- context : system context, address (memory) context, H/W context

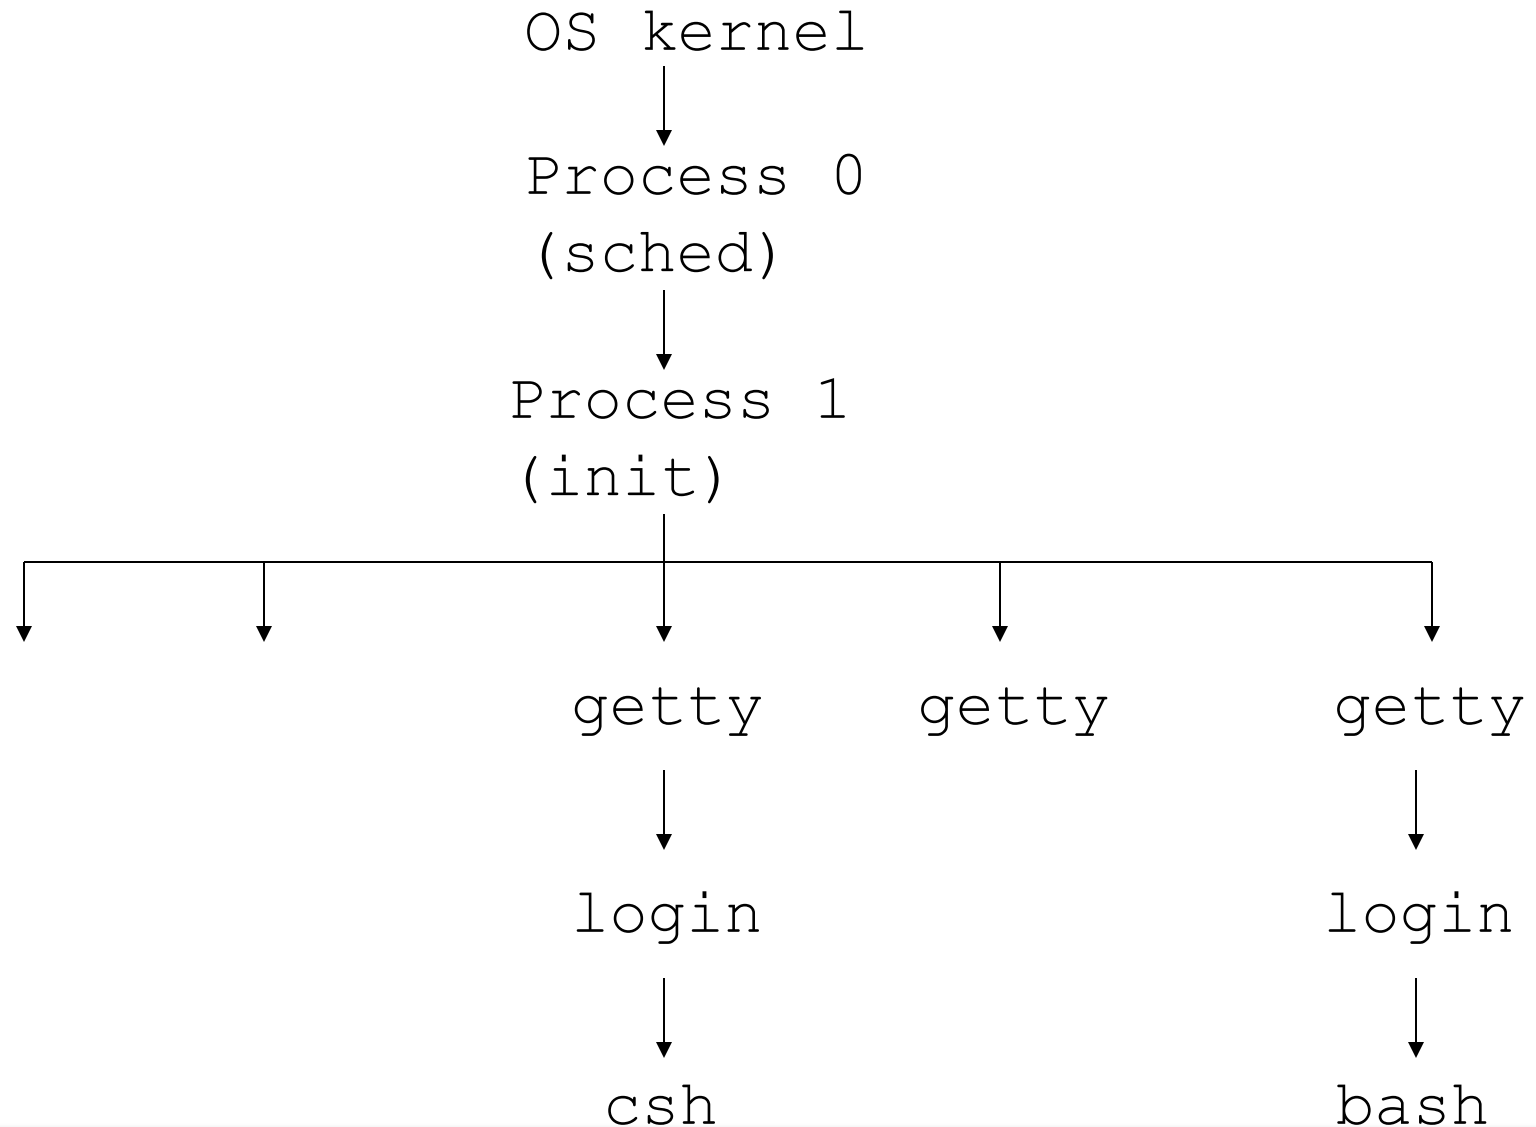


# Important System Processes

- **init** – Mother of all processes. init is started at boot time and is responsible for starting other processes.
  - init uses file inittab & directories: /etc/rc?.d
- **getty** – login process that manages login sessions.



# Unix Start Up Processes Diagram





# Pid and Parentage

- A process ID or *pid* is a positive integer that uniquely identifies a running process, and is stored in a variable of type *pid\_t*.
- You can get the **process pid** or **parent's pid**

```
#include <sys/types>
main()
{
    pid_t pid, ppid;
    printf( "My  PID is:%d\n\n", (pid = getpid()) );
    printf( "Par PID is:%d\n\n", (ppid = getppid()) );
}
```



## 2. fork()

- ```
#include <sys/types.h>
#include <unistd.h>
pid_t fork( void );
```
- Creates a child process by making a copy of the parent process --- an **exact** duplicate.
  - Implicitly specifies code, registers, stack, data, files
- Both the child *and* the parent continue running.



# Process IDs (pids revisited)

- `pid = fork();`
- In the child: `pid == 0;`  
In the parent: `pid ==` the process ID of the child.
- A program almost always uses this `pid` difference to do different things in the parent and child.



# fork() Example (parchld.c)

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main()
{
    pid_t pid;           /* could be int */
    int i;
    pid = fork();
    if( pid > 0 )
    {
        /* parent */
        for( i=0; i < 1000; i++ )
            printf("\t\t\tPARENT %d\n", i);
    }
}
```



```
else
{
    /* child */
    for( i=0; i < 1000; i++ )
        printf( "CHILD %d\n", i );
    }
    return 0;
}
```



# Possible Output

CHILD 0

CHILD 1

CHILD 2

PARENT 0

PARENT 1

PARENT 2

PARENT 3

CHILD 3

CHILD 4

PARENT 4

:



# Things to Note

- $i$  is copied between parent and child.
- The switching between the parent and child depends on many factors:
  - machine load, system process scheduling
- Output interleaving is *nondeterministic*
  - cannot determine output by looking at code



# LAB7-1: Address context

---

## ■ *fork example*

```
int glob = 6;
char buf[] = "a write to stdout\n";

int main(void)
{
    int var;
    pid_t pid;

    var = 88;
    write(STDOUT_FILENO, buf, sizeof(buf)-1);
    printf("before fork\n");

    if ((pid = fork()) == 0) {    /* child */
        glob++; var++;
    } else
        sleep(2);               /* parent */

    printf("pid = %d, glob = %d, var = %d\n", getpid(), glob, var);
    exit (0);
}
```

(Source : Adv. programming in the UNIX Env., pgm 8.1)

👉 guess what can we get from this program?






## LAB7-1: Address context

- Since a child process is a **copy** of the parent, it has copies of the parent's data.
- A change to a variable in the child will *not* change that variable in the parent.



# LAB7-1: Address context

- fork internal : **before fork** (after run a.out)

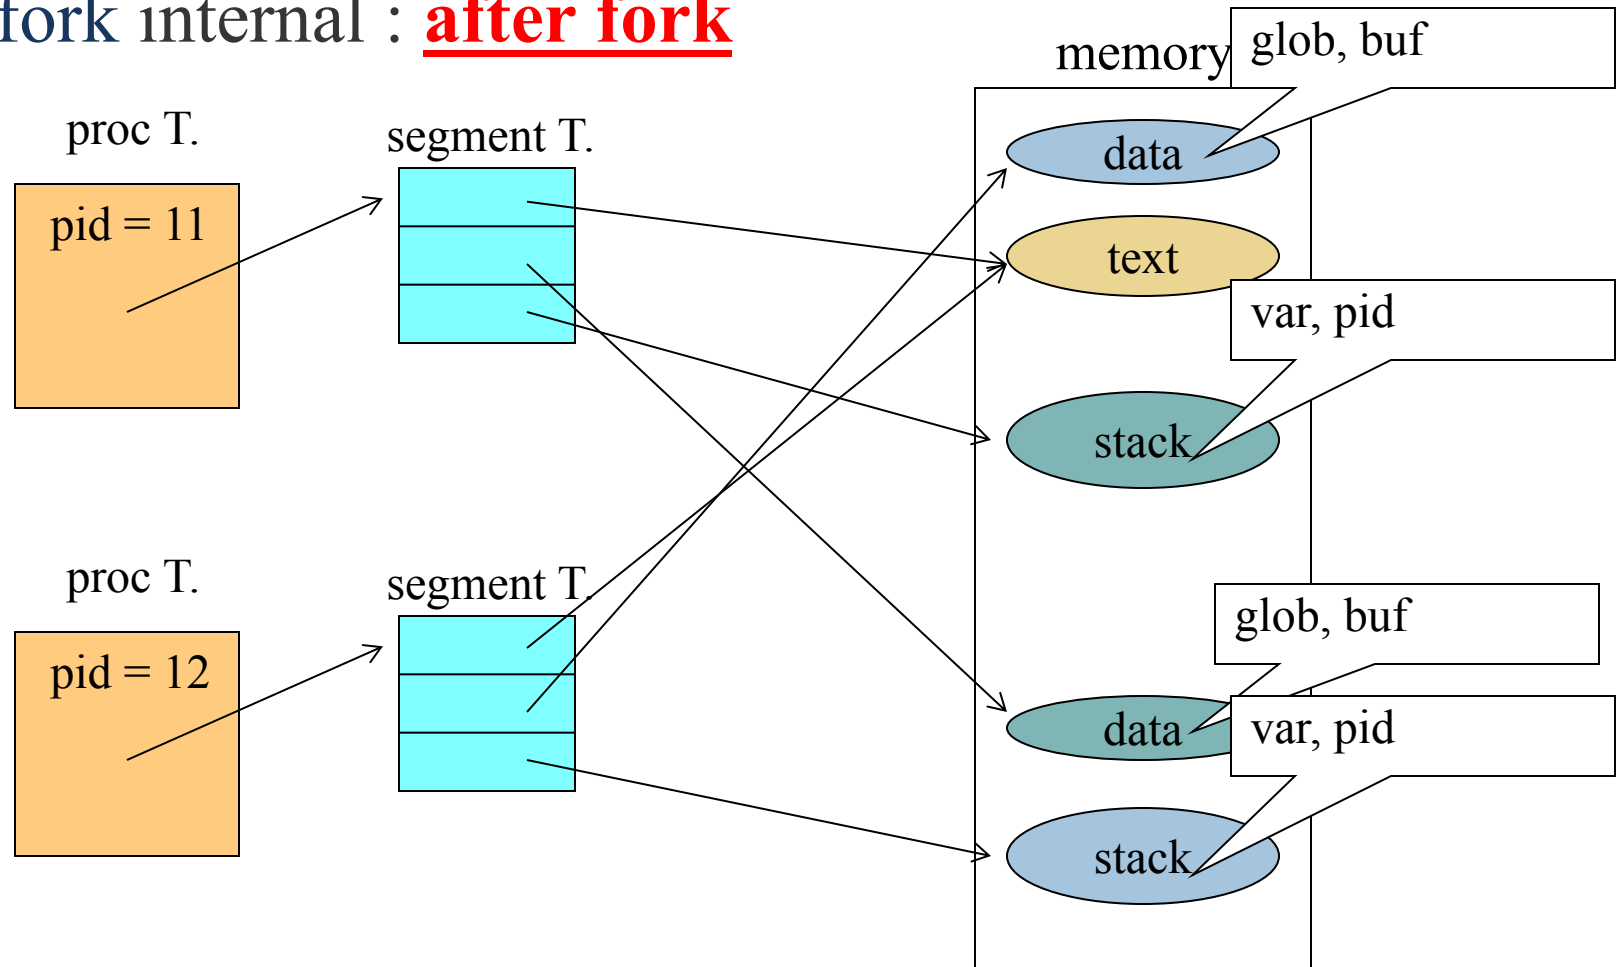


We assume that there is no paging mechanism in this figure.



# LAB7-1: Address context

- fork internal : **after fork**



# LAB7-1: Address context

## ■ Answer

- *Note: we never know which of the parent and the child starts first.*

```
$ ./a.out  
a write to stdout  
before fork
```

Only once. Write() is **NOT** buffered.

Only once. printf() is buffered, but flushed by newline.

```
pid = 430, glob = 7, var = 89    child's variables were changed  
pid = 429, glob = 6, var = 88    parent's copy was not changed  
$ ./a.out > temp.out  
$ cat temp.out  
a write to stdout  
before fork  
pid = 432, glob = 7, var = 89  
before fork  
pid = 431, glob = 6, var = 88
```

Twice. printf() is buffered and **NOT** flushed. Copied to the child.



### 3. exec()

- Family of functions for **replacing** process's program with the one inside the `exec()` call.

```
#include <unistd.h>

int execl(const char *pathname, const char *arg0, ... /* (char *)0 */ );
int execv(const char *pathname, char *const argv[]);
int execl(const char *pathname, const char *arg0, ...
          /* (char *)0, char *const envp[] */ );
int execve(const char *pathname, char *const argv[], char *const envp[]);
int execlp(const char *filename, const char *arg0, ... /* (char *)0 */ );
int execvp(const char *filename, char *const argv[]);
int fexecve(int fd, char *const argv[], char *const envp[]);
```

All seven return: -1 on error, no return on success



# exec(..) Family

- There are 6 versions of the exec function, and they all do about the same thing: they replace the current program with the text of the new program. Main difference is how parameters are passed.



```
int execl( const char *path, const char *arg, ... );
int execlp( const char *file, const char *arg, ... );
int execl( const char *path, const char *arg
           , ..., char *const envp[] );
int execv( const char *path, char *const argv[] );
int execvp( const char *file, char *const argv[] );
int execve( const char *filename, char *const argv [],
            char *const envp[] );
```



### 3. exec()

- Example

```
#include <unistd.h>

int execlp(char *file, char *arg0,
           char *arg1, ..., (char *)0);

execlp("sort", "sort", "-n",
       "foobar", (char *)0);
```

↑  
Same as "sort -n foobar"





## LAB7-2: tinymenu.c


```
#include <stdio.h>
#include <unistd.h>

void main()
{
    char *cmd[] = {"who", "ls", "date"};
    int i;
    printf("0=who 1=ls 2=date : ");
    scanf("%d", &i);

    execvp( cmd[i], cmd[i], (char *)0 );
    printf( "execvp failed\n" );
}
```

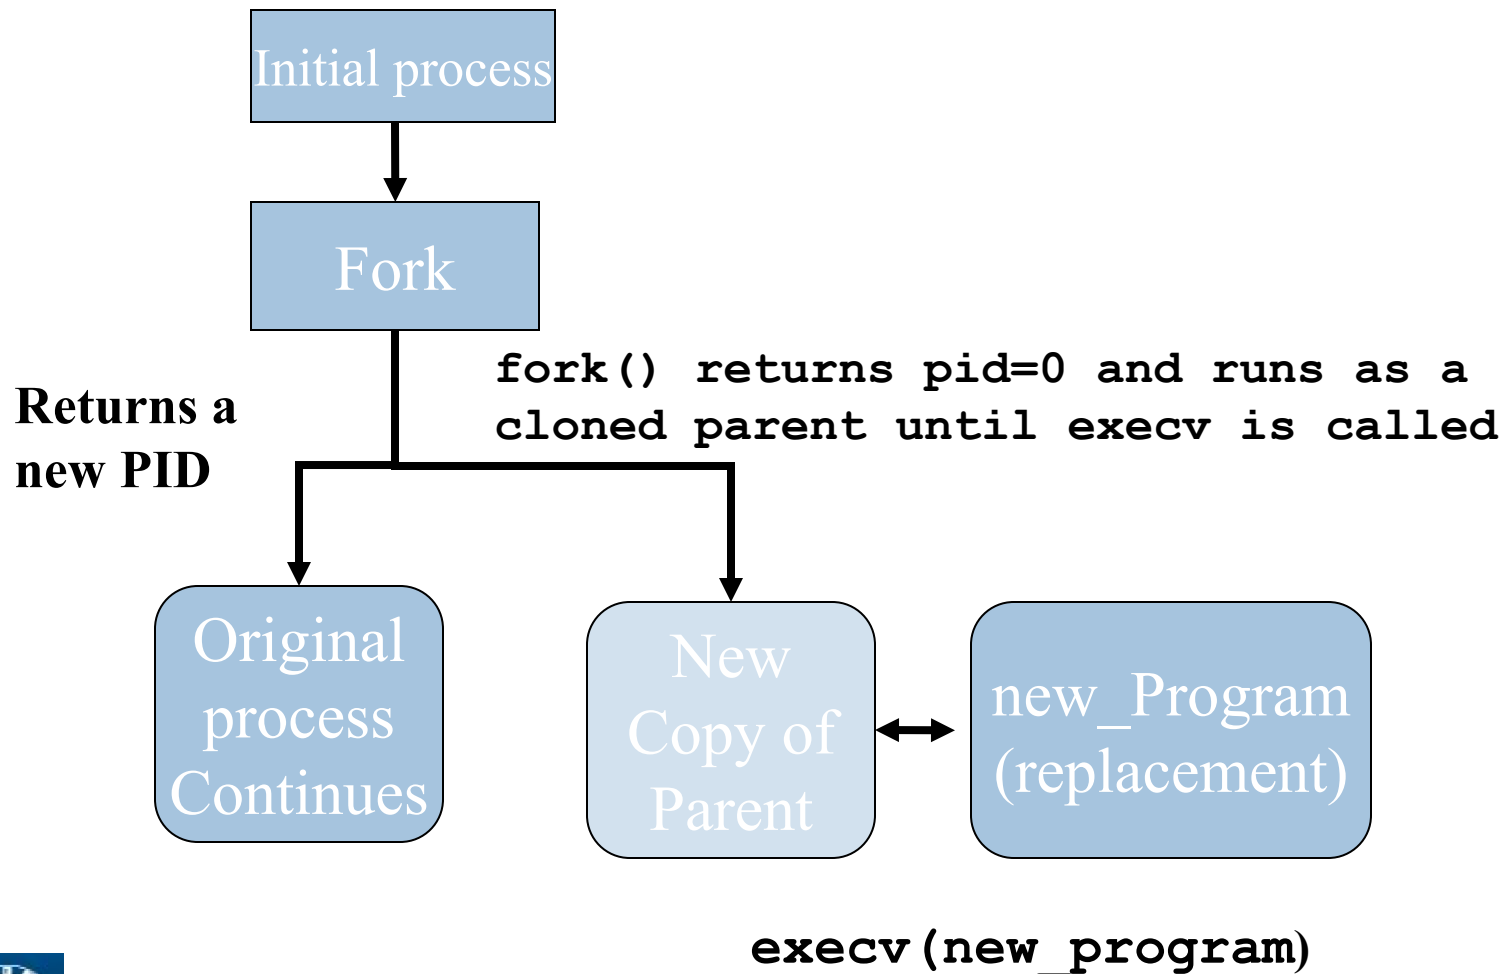


# Execution



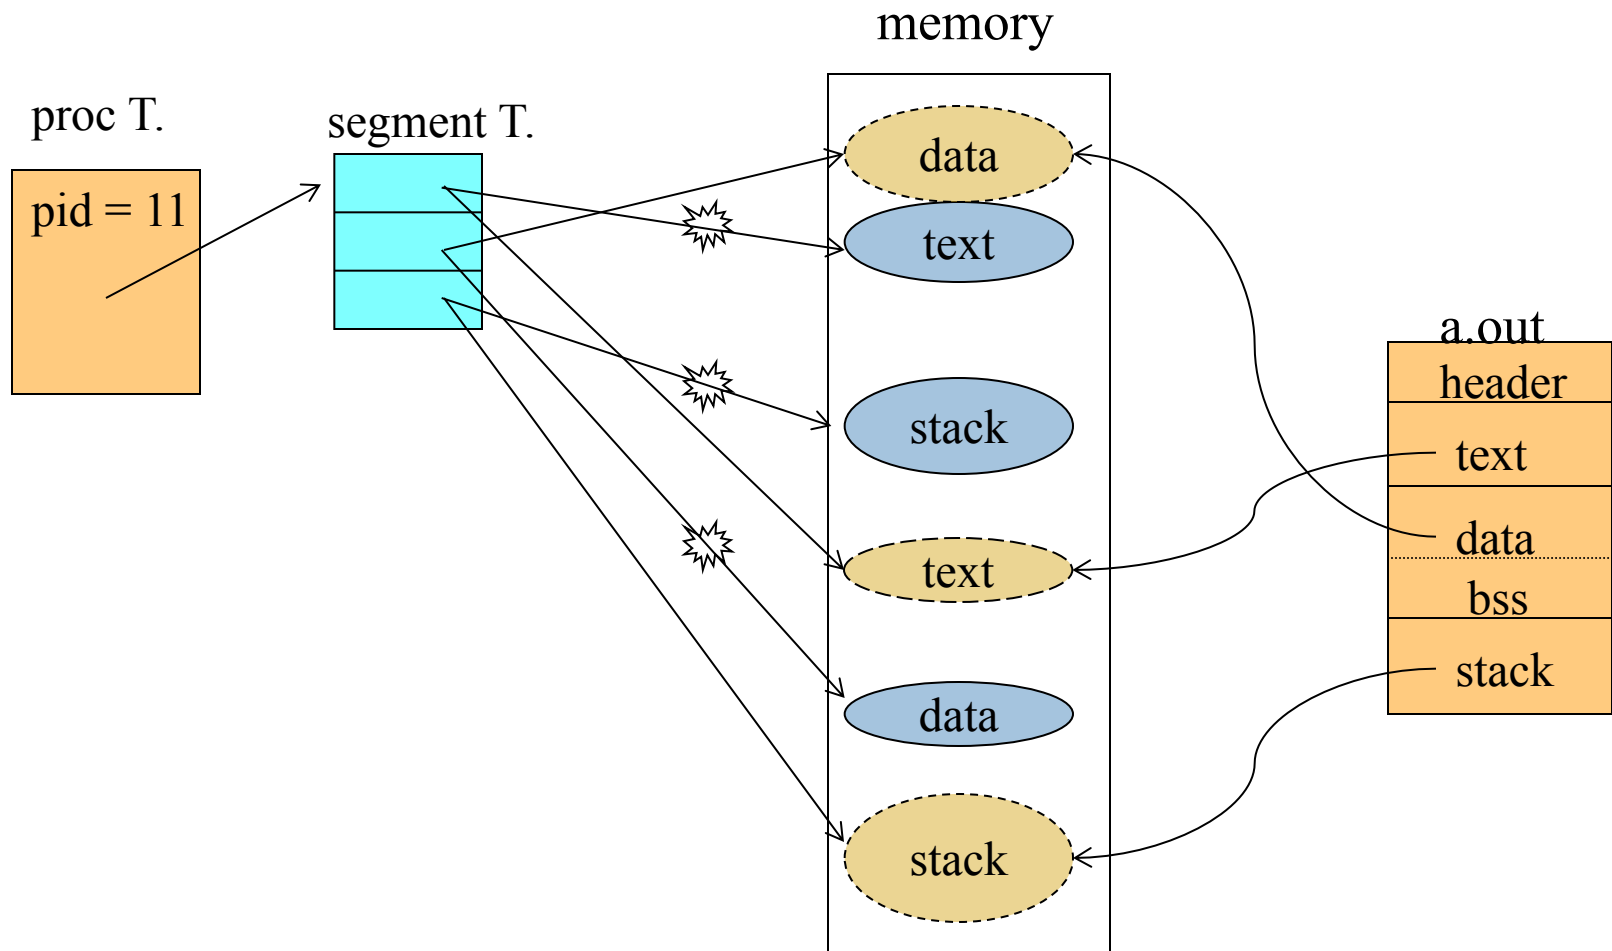
# fork() and execv()

- `execv(new_program, argv[ ])`



# Context : address context

- execve internal



## 4. wait()

- ```
#include <sys/types.h>
#include <sys/wait.h>
pid_t wait(int *statloc);
```
- **Suspends calling process until child has finished.**  
Returns the process ID of the terminated child if ok, -1 on error.
- `statloc` can be `(int *)0` or a variable which will be bound to status info. about the child.



## `wait()` Actions

- A process that calls `wait()` can:
  - *suspend* (block) if all of its children are still running, or
  - *return* immediately with the *termination* status of a child, or
  - *return* immediately with an *error* if there are no child processes.



## LAB7-3: menushell.c

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>

void main()
{
    char *cmd[] = {"who", "ls", "date"};
    int i;
    while( 1 )
    {
        printf( 0=who 1=ls 2=date : " );
        scanf( "%d", &i );
        :
```

*continued*

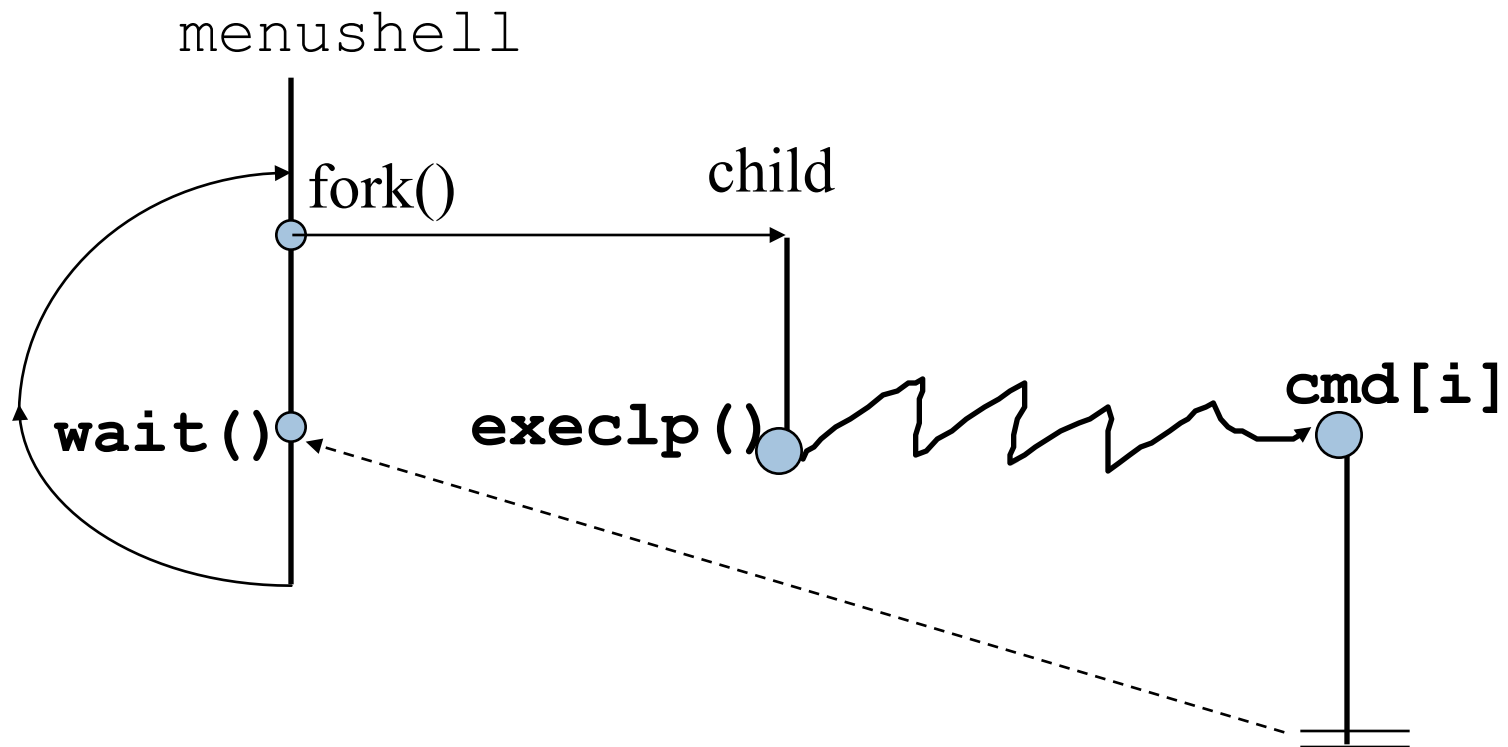


```
if(fork() == 0)
    { /* child */
        execlp( cmd[i], cmd[i], (char *)0 );
        printf( "execlp failed\n" );
        exit(1);
    }
else
    { /* parent */
        wait( (int *)0 );
        printf( "child finished\n" );
    }
} /* while */
} /* main */
```





# Execution



# Macros for wait (1)

- **WIFEXITED(*status*)**
  - Returns true if the child exited normally.
- **WEXITSTATUS(*status*)**
  - Evaluates to *the least significant eight bits* of the return code of the child which terminated, which may have been set as the argument to a call to `exit( )` or as the argument for a return.
  - This macro can only be evaluated if **WIFEXITED** returned non-zero.



## Macros for wait (2)

- **WIFSIGNALED(*status*)**
  - Returns true if the child process exited *because of a signal* which was not caught.
- **WTERMSIG(*status*)**
  - Returns *the signal number* that caused the child process to terminate.
  - This macro can only be evaluated if **WIFSIGNALED** returned non-zero.



# waitpid()

```
#include <sys/types.h>
#include <sys/wait.h>
pid_t waitpid( pid_t pid, int *status, int opts )
```

- waitpid - can wait for a particular child
- *pid* < -1
  - Wait for **any child process whose process group ID** is equal to the absolute value of *pid*.
- *pid* == -1
  - Wait for **any child process**.
  - Same behavior which wait( ) exhibits.
- *pid* == 0
  - Wait for **any child process whose process group ID** is equal to that of the calling process.



- *pid* > 0
  - Wait for the child whose process ID is equal to the value of *pid*.
  - *options*
    - Zero or more of the following constants can be ORed.
      - WNOHANG
        - » Return immediately if no child has exited.
      - WUNTRACED
        - » Also return for children which are stopped, and whose status has not been reported (because of signal).
  - Return value
    - The process ID of the child which exited.
    - -1 on error; 0 if WNOHANG was used and no child was available.



# Macros for waitpid

- WIFSTOPPED(*status*)
  - Returns true if the child process which caused the return is *currently stopped*.
  - This is only possible if the call was done using WUNTRACED.
- WSTOPSIG(*status*)
  - Returns *the signal number* which caused the child to stop.
  - This macro can only be evaluated if WIFSTOPPED returned non-zero.



# LAB7-4: waitpid

```
#include <stdio.h>
#include <sys/wait.h>
#include <sys/types.h>

int main(void)
{
    pid_t pid;
    int status;

    if( (pid = fork() ) == 0 )
    { /* child */
        printf("I am a child with pid = %d\n",
            getpid());
        sleep(60);
        printf("child terminates\n");
        exit(0);
    }
```



```
else
{ /* parent */
    while (1)
    {
        waitpid( pid, &status, WUNTRACED );
        if( WIFSTOPPED(status) )
        {
            printf("child stopped, signal(%d)\n",
                    WSTOPSIG(status));
            continue;
        }
        else if( WIFEXITED(status) )
            printf("normal termination with
                    status(%d)\n",
                    WEXITSTATUS(status));
        else if (WIFSIGNALED(status))
            printf("abnormal termination,
                    signal(%d)\n",
                    WTERMSIG(status));

        exit(0);
    } /* while */
} /* parent */
} /* main */
```





## 5. Process File Descriptors

- A child and parent have **copies** of the file descriptors, **but the R-W pointer is maintained by the system:**
  - the R-W pointer is shared.
- This means that a `read()` or `write()` in one process will affect the other process since the R-W pointer is changed.



## LAB7-5: File used across processes (shfile.c)

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <fcntl.h>
void printpos(char *msg, int fd);
void fatal(char *msg);

int main(void)
{
    int fd;      /* file descriptor */
    pid_t pid;
    char buf[10]; /* for file data */
    :
```

*continued*



```
if ( (fd=open("data-file", O_RDONLY)) < 0 )
    perror("open");

read(fd, buf, 10); /* move R-W ptr */

printpos( "Before fork", fd );

if( (pid = fork()) == 0 )
{ /* child */
    printpos( "Child before read", fd );
    read( fd, buf, 10 );
    printpos( " Child after read", fd );
}

:
```

*continued*



```
else if( pid > 0 )
{
    /* parent */
    wait((int *)0);
    printpos( "Parent after wait", fd );
}
else
    perror( "fork" );
}
```

*continued*



```
void printpos( char *msg, int fd )
/* Print position in file */
{
    long int pos;

    if( (pos = lseek( fd, 0L, SEEK_CUR) ) < 0L )
        perror("lseek");

    printf( "%s: %ld\n", msg, pos );
}
```



# Output

```
$ shfile
```

```
Before fork: 10
```

```
Child before read: 10
```

```
Child after read: 20
```

```
Parent after wait: 20
```

what's happened?



## 7. Special Exit Cases

Two special cases:

1. A child exits when its parent is not currently executing `wait()`
  - the child becomes a *zombie*
  - `status` data about the child is stored until the parent does a `wait()`



2. A parent exits when 1 or more children are still running

- children are **adopted** by the system's initialization process (`/etc/init`)
  - it can then monitor/kill them








## 8. I/O redirection

- The trick: you can change where the standard I/O streams are going/coming from **after** the fork **but** before the exec



# Redirection of standard output

- Example implement shell: **ls > x.lis**
- program:
  - Open a new file x.lis
  - Redirect standard output to x.lis using **dup** command
    - everything sent to standard output ends in x.lis
  - execute **ls** in the process
- **dup2(int fin, int fout)** - copies fin to fout in the fd table



# LAB: implement ls > x.lis

```
#include <unistd.h>
int main ()
{
    int fileId;
    fileId = creat( "x.lis", 0640 );
    if( fileId < 0 )
    {
        printf( stderr, "error creating x.lis\n" );
        exit (1);
    }
    dup2( fileId, stdout ); /* copy fileId to stdout */
    close( fileId );
    execl( "/bin/ls", "ls", 0 );
}
```



# dup, dup2

```
#include <unistd.h>
```

```
int dup(int filedes);
```

```
int dup2(int filedes, int filedes2);
```


Both return : new file descriptor if OK, -1 on error

- An existing file descriptor (`filedes`) is duplicated
- The new file descriptor returned by `dup` is guaranteed to be the lowest numbered available file descriptor
- With `dup2`, we specify the value of the new descriptor with the `filedes2` argument
- If `filedes2` is already open, it is first closed.
- Ex) `dup2 (fd, STDOUT_FILENO) ;`



# example

- `dup(1);`



# LAB: What does this program do?



```
#include <fcntl.h>           // redirection.c
#include <unistd.h>
#include <stdio.h>
#include <errno.h>
#include <sys/wait.h>

int main(int argc , char* argv[]) {
    char *path = "/bin/ls";
    char *arg0 = "ls";
    pid_t pid;
    int fd;
    int status;

    if (argc!=2) {
        printf("wrong command. ex) a.out filename\n");
        exit(0);
    }
    pid = fork();
    if (pid == 0) {
        fd = open(argv[1], O_WRONLY|O_CREAT|O_TRUNC, S_IRWXU);
        dup2(fd, STDOUT_FILENO);
        close(fd);
        if (execl(path, arg0, NULL) == -1)
            perror("execl");
    } else {
        close(fd);
        wait(&status);
    }
}
```

# Pipe

```
$ who | sort
```





# pipe

```
#include <unistd.h>

int pipe(int filedes[2]);

        return : 0 if OK, -1 on error
```



Two file descriptors are **returned** through the *fd* argument

- *fd*[0]: can be used to **read** from the pipe, and
- *fd*[1]: can be used to **write** to the pipe



**Anything that is written on *fd*[1] may be read by *fd*[0].**

- This is **of no use in a single process.**
- However, between processes, it gives a method of communication



The pipe() system call gives parent-child processes a way to communicate with each other.





# pipe

- `filedes[0]` is open for reading and `filedes[1]` is open for writing.
- Output of `filedes[1]` is the input for `filedes[0]`



# pipe after fork




# Pipes

**parent → child:**

parent closes fd[0]


child closes fd[1]



**parent ← child:**

parent closes fd[1]

child closes fd[0]



# LAB: Pipe example results?

```
#include <stdio.h> // pipe2.c

#define READ  0
#define WRITE 1

char* phrase = "Stuff this in your pipe and smoke it";

main( ) {
    int fd[2], bytesRead;
    char message[100];
    pipe(fd);
    if (fork() == 0) { // child
        close(fd[READ]);
        write(fd[WRITE], phrase, strlen(phrase)+1);
        fprintf(stdout, "[%d, child] write completed.\n", getpid());
        close(fd[WRITE]);
    }
    else { // parent
        close(fd[WRITE]);
        bytesRead = read(fd[READ], message, 100);
        fprintf(stdout, "[%d, parent] read completed.\n", getpid());
        printf("Read %d bytes: %s\n", bytesRead, message);
        close(fd[READ]);
    }
}
```



## Another Example

```
#include <sys/types.h>    // pipe.c
#include <unistd.h>
#include <stdio.h>
#include <errno.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    char *path = "/bin/ls";
    char *arg0 = "ls";
    pid_t pid;
    int pipefd[2];
    int status;
    pipe(pipefd);
    pid = fork();
    if (pid == 0) {
        dup2(pipefd[1], STDOUT_FILENO);
        close(pipefd[0]);
        close(pipefd[1]);
        if (execl(path, arg0, NULL) == -1)
            perror("execl");
    } else {
        if (fork() == 0) {
            dup2(pipefd[0], STDIN_FILENO);
            close(pipefd[0]);
            close(pipefd[1]);
            if (execl("/bin/cat", "cat", NULL) == -1)
                perror("execl cat");
        } else {
            close(pipefd[0]);
            close(pipefd[1]);
            wait(&status);
            wait(&status);
        }
    }
}
```