



Process Environment (Chap7-Chap9)

Hongik University

Eunsung Jung

Disclaimer: The slides are borrowed from many sources!

LAB: Shell Command-Line Processing

- Refer to Figure 7.4

```
$ cc -Wall argv.c
$ ./a.out
$ ./a.out *.c
$ ./a.out *.none
$ ./a.out *. [1c]
$ ./a.out "*.c"
$ ./a.out $USER
$ ./a.out "$(echo *.1)"
$ ./a.out {foo,bar,baz}.whatever
$ ./a.out {1..5}
$ ./a.out {1..5}{a..f}
```

See also: <https://is.gd/kJXbpa> and <http://is.gd/iZa9rC>



Memory Layout of a C Program

- The `size (1)` command reports the sizes (in bytes) of the text, data, and bss segments.

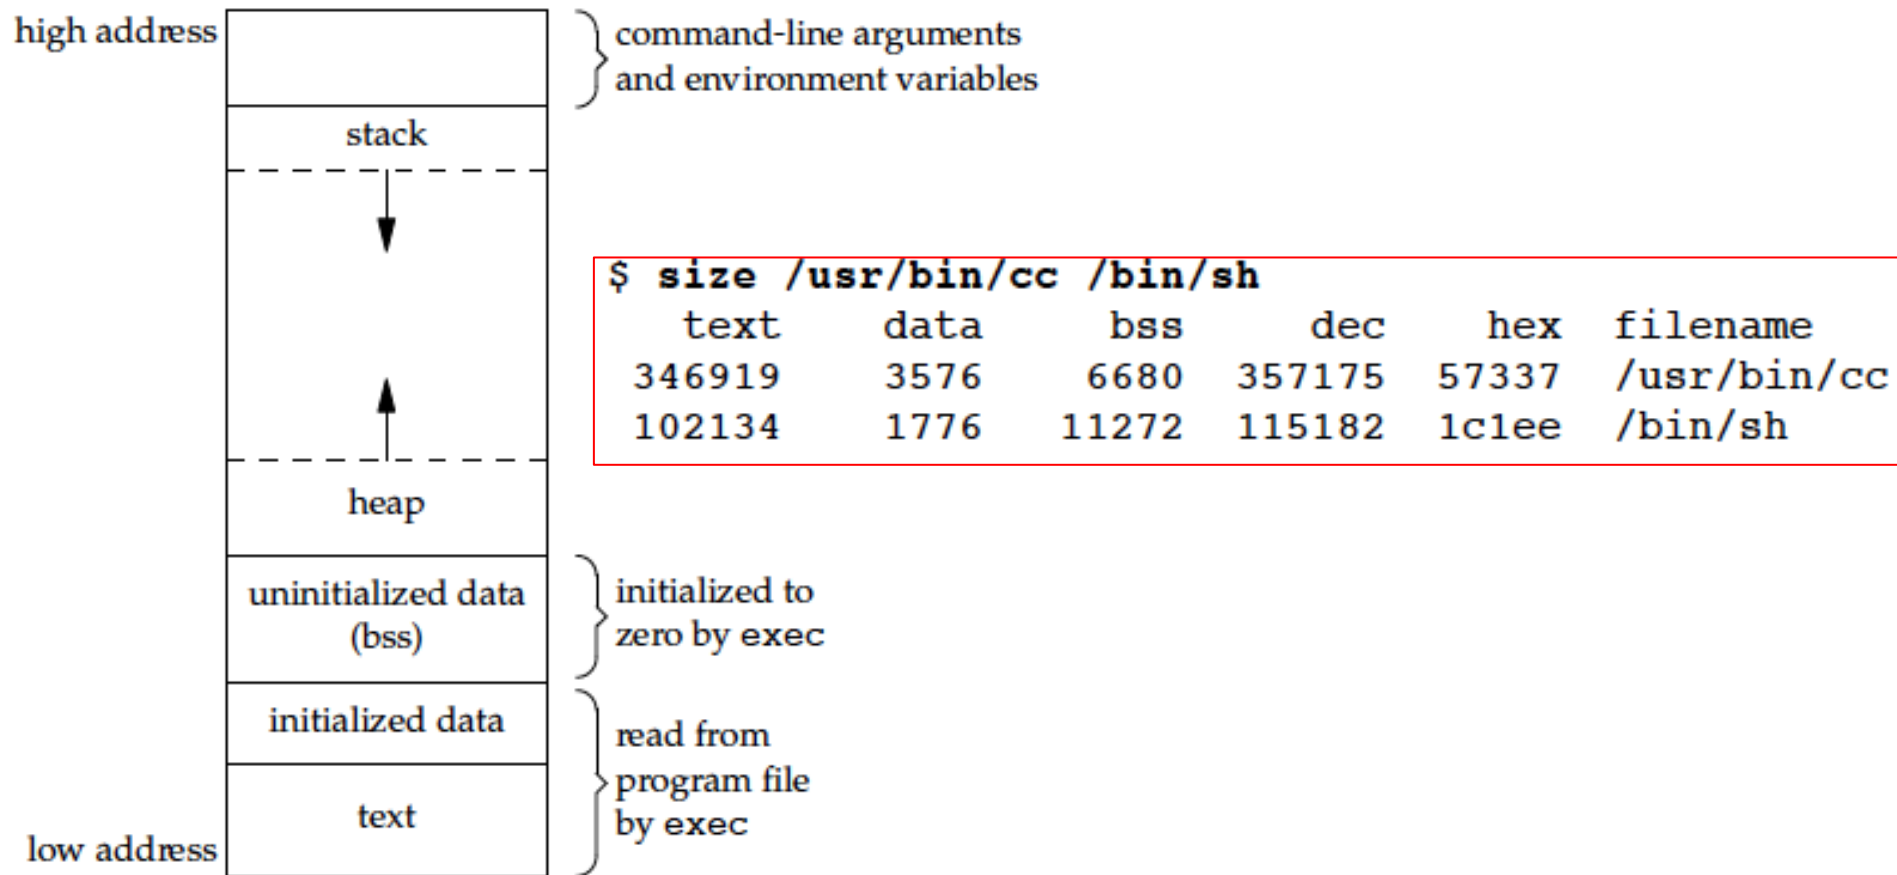


Figure 7.6 Typical memory arrangement

The main function


```
int main(int argc, char **argv);
```

- C program started by **kernel** (by one of the exec functions)
- **Special startup routine** called by kernel which sets up things for main (or whatever entry point is defined)
- `argc` is a count of the number of command line arguments (including the command itself)
- `argv` is an array of pointers to the arguments
- it is guaranteed by both ANSI C and POSIX.1 that `argv[argc] == NULL`



Process Creation

- **Executable and Linkable Format (ELF, formerly named Extensible Linking Format)**



Process Creation

- On Linux:

```
$ cc -Wall entry.c
$ readelf -h a.out | more
```

```
ELF Header:
```

```
[...]
```

Entry point address:	0x400460
Start of program headers:	64 (bytes into file)
Start of section headers:	4432 (bytes into file)

```
$ objdump -d a.out
```

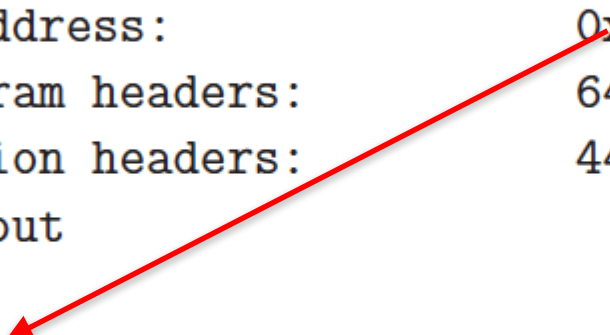
```
[...]
```

```
0000000000400460 <_start>:
```

400460:	31 ed	xor	%ebp,%ebp
400462:	49 89 d1	mov	%rdx,%r9

```
[...]
```

```
$
```



Process Creation

glibc/sysdeps/x86_64/start.S

0000000000401058 <_start>:

401058:	31 ed	xor	%ebp,%ebp
40105a:	49 89 d1	mov	%rdx,%r9
40105d:	5e	pop	%rsi
40105e:	48 89 e2	mov	%rsp,%rdx
401061:	48 83 e4 f0	and	\$0xfffffffffffffffff0,%rsp
401065:	50	push	%rax
401066:	54	push	%rsp
401067:	49 c7 c0 e0 1a 40 00	mov	\$0x401ae0,%r8
40106e:	48 c7 c1 50 1a 40 00	mov	\$0x401a50,%rcx
401075:	48 c7 c7 91 11 40 00	mov	\$0x401191,%rdi
40107c:	e8 2f 01 00 00	callq	4011b0 <__libc_start_main>
401081:	f4	hlt	
401082:	90	nop	
401083:	90	nop	



Process Creating

- `git clone git://sourceware.org/git/glibc.git`
- `cd glibc`
- Find `start.S`!
- FYI, Kernel codes
 - `wget https://www.kernel.org/pub/linux/kernel/v4.x/linux-4.12.12.tar.gz`
 - `tar xvf linux-4.12.12.tar.gz`



Process Creation

glibc/csu/libc-start.c

STATIC int

```
LIBC_START_MAIN (int (*main) (int, char **, char ** MAIN_AUXVEC_DECL),  
                int argc, char **argv,  
                __typeof (main) init,  
                void (*fini) (void),  
                void (*rtld_fini) (void), void *stack_end)  
{  
    [...]  
    result = main (argc, argv, __environ MAIN_AUXVEC_PARAM);  
  
    exit (result);  
}
```



Process Termination

- There are 8 ways for a process to terminate.
- Normal termination (5 ways)
 - return from main
 - calling `exit`
 - calling `_exit` (or `_Exit`)
 - return of last thread from its start routine
 - calling `pthread_exit` from last thread



Process Termination

- There are 8 ways for a process to terminate.
- Abnormal termination (3 ways)
 - calling abort
 - terminated by a signal
 - response of the last thread to a cancellation request



exit(3) and _exit(2)

```
#include <stdlib.h>

void exit(int status);
void _Exit(int status);

#include <unistd.h>
void _exit(int status);
```

- `_exit` and `_Exit`
 - return to the kernel immediately
 - `_exit` required by POSIX.1
 - `_Exit` required by ISO C99
 - synonymous on Unix
- `exit` does some cleanup and then returns
- both take integer argument, aka exit status



atexit(3)

```
#include <stdlib.h>

int atexit(void (*func)(void));
```

- Registers a function with a signature of void funcname(void) to be called at exit
- Functions invoked in reverse order of registration
- Same function can be registered more than once
- Extremely useful for cleaning up open files, freeing certain resources, etc.



Lifetime of a UNIX Process

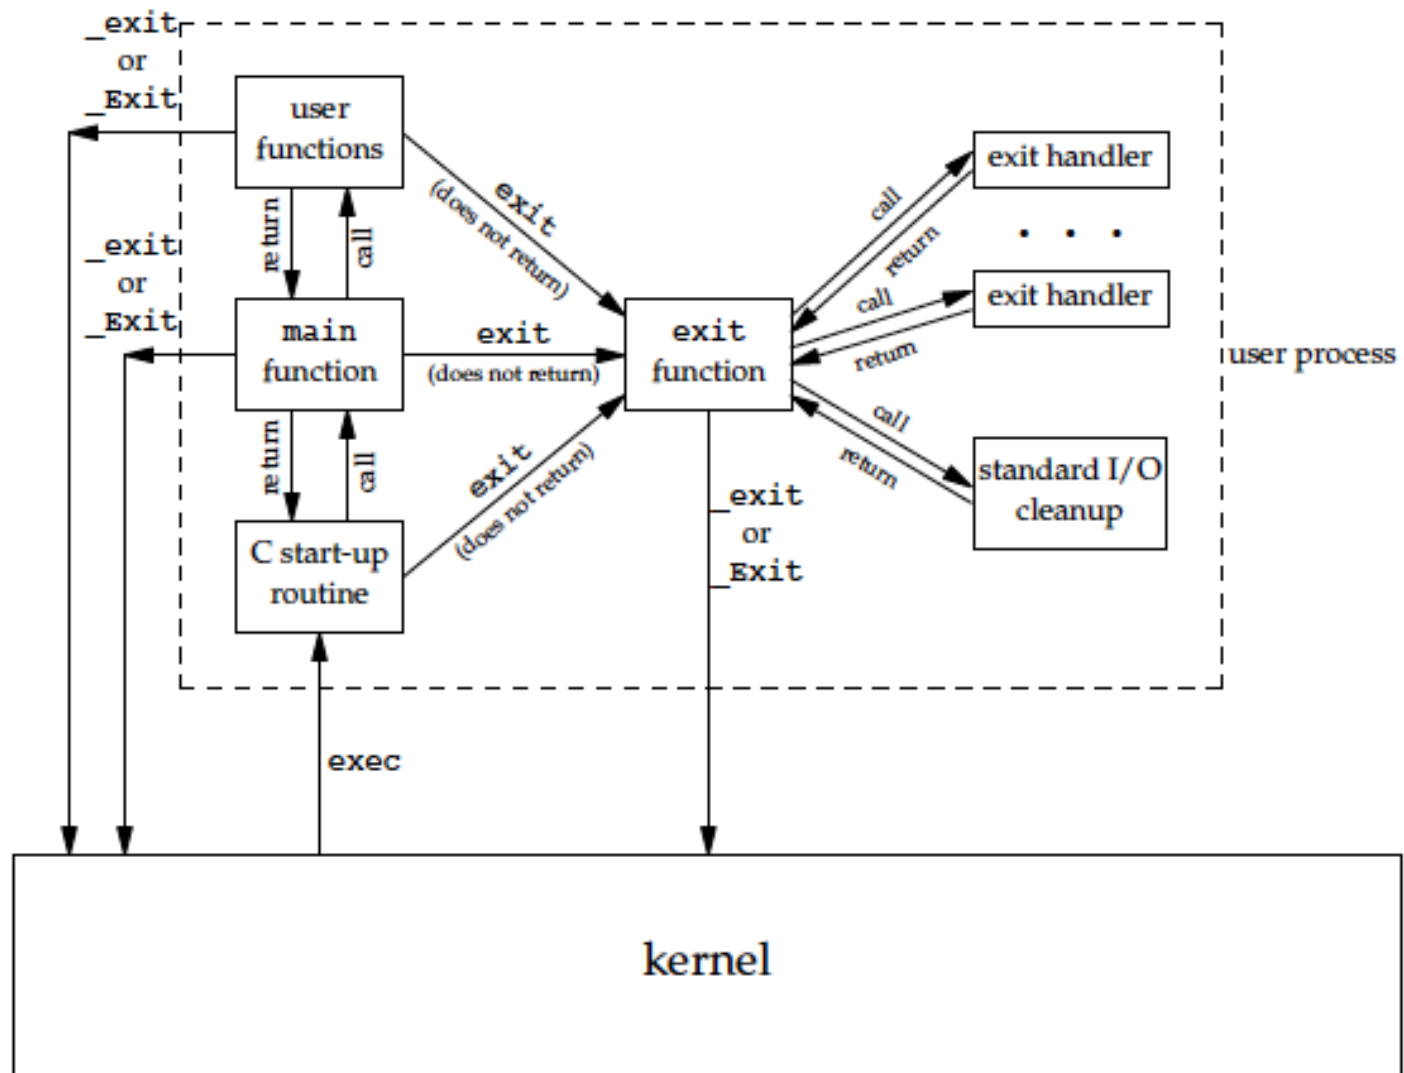


Figure 7.2 How a C program is started and how it terminates



Environment List

- Environment variables are stored in a global array of pointers:
 - extern char **environ;
 - The list is null terminated.
- These can also be accessed by:

```
#include <stdlib.h>

char *getenv(const char *name);
int putenv(const char *string);
int setenv(const char *name, const char *value, int rewrite);
void unsetenv(const char *name);
```

```
int main(int argc, char **argv, char **envp);
```



Memory Allocation

```
#include <stdlib.h>

void *malloc(size_t size);
void *calloc(size_t nobj, size_t size);
void *realloc(void *ptr, size_t newsiz);
void *alloca(size_t size);

void free(void *ptr);
```

- **malloc** – initial value is indeterminate.
- **calloc** – initial value set to all zeros.
- **realloc** – changes size of previously allocated area. Initial value of any additional space is indeterminate.
- **alloca** – allocates memory on stack



Memory Layout of a C Program

On NetBSD:

```
$ cc hw.c
```

```
$ file a.out
```

```
a.out: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically  
linked (uses shared libs), for NetBSD 5.0, not stripped
```

```
$ ldd a.out
```

```
a.out:
```

```
    -lc.12 => /usr/lib/libc.so.12
```

```
$ size a.out
```

text	data	bss	dec	hex filename
2301	552	120	2973	b9d a.out

```
$ objdump -d a.out > obj
```

```
$ wc -l obj
```

```
271 obj
```

```
$
```



Memory Layout of a C Program

On Mac OS X:

```
$ cc hw.c
```

```
$ file a.out
```

```
a.out: Mach-O 64-bit executable x86_64
```

```
$ otool -L a.out
```

```
a.out:
```

```
/usr/lib/libSystem.B.dylib (compatibility version 1.0.0,  
current version 125.2.11)
```

```
$ size a.out
```

```
__TEXT __DATA __OBJC others dec hex  
4096 4096 0 4294971392 4294979584 100003000
```

```
$ otool -t -v a.out > obj
```

```
$ wc -l obj  
32 obj
```

```
$
```



Memory Layout of a C Program

On Linux:

```
$ cc hw.c
$ file a.out
a.out: ELF 32-bit LSB executable, Intel 80386, version 1 (SYSV),
dynamically linked (uses shared libs), for GNU/Linux 2.6.15, not stripped
$ ldd a.out
linux-gate.so.1 => (0x00c66000)
libc.so.6 => /lib/tls/i686/cmov/libc.so.6 (0x006b4000)
/lib/ld-linux.so.2 (0x005fe000)
$ size a.out
   text    data     bss      dec     hex filename
   918     264        8    1190    4a6 a.out
$ objdump -d a.out >obj
$ wc -l obj
225 obj
$
```



Process limits

```
$ ulimit -a
time(cpu-seconds)      unlimited
file(blocks)           unlimited
coredump(blocks)       unlimited
data(kbytes)           262144
stack(kbytes)          2048
lockedmem(kbytes)      249913
memory(kbytes)         749740
nofiles(descriptors)   128
processes              160
vmemory(kbytes)        unlimited
sbsize(bytes)          unlimited
$
```



getrlimit(2) and setrlimit(2)

```
#include <sys/resource.h>

int getrlimit(int resource, struct rlimit *rlp);
int setrlimit(int resource, const struct rlimit *rlp);
```

- Changing resource limits follows these rules:
 - a **soft limit can be changed** by any process to a value less than or equal to its **hard limit**
 - any process can lower its hard limit greater than or equal to its soft limit
 - only **superuser** can raise hard limits
 - changes are per process only

