



Chap 10. Signals

Hongik University

Eunsung Jung

Disclaimer: The slides are borrowed from many sources!

Signals

- Signals are **software** interrupts from unexpected events
 - an illegal operation (e.g., divide by 0)
 - a power failure
 - an alarm clock
 - the death of a child process
 - a termination request from a user (Ctrl-C)
 - a suspend request from a user (Ctrl-Z)

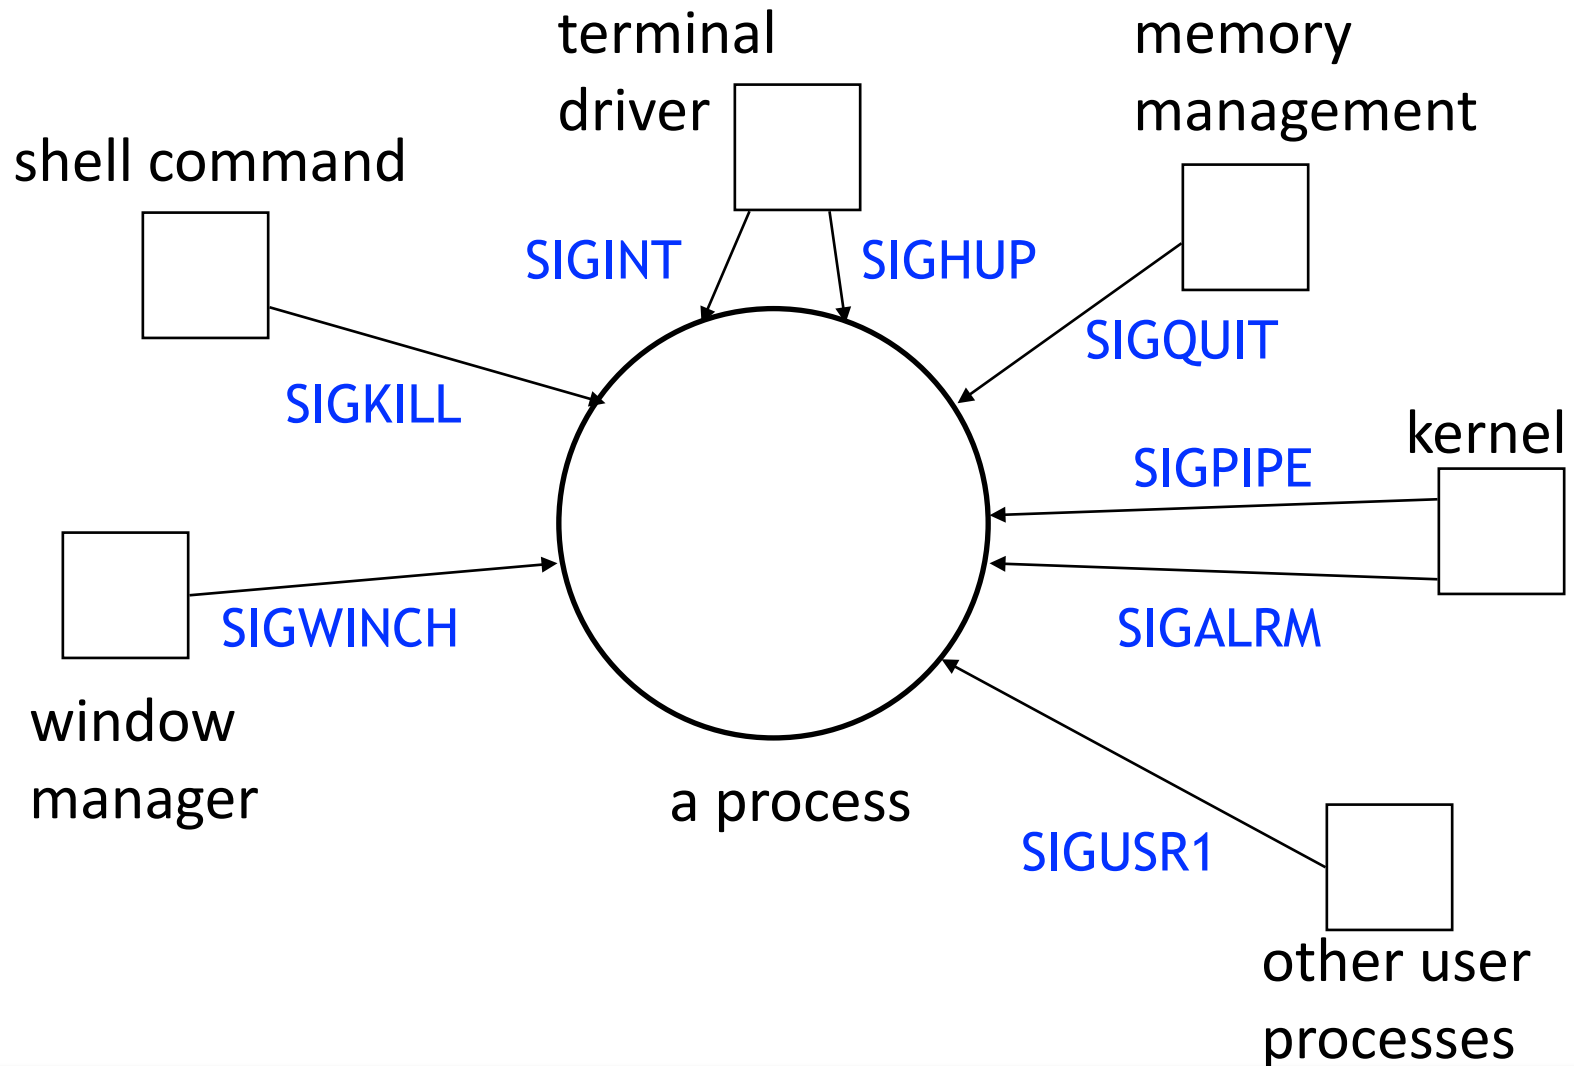


Predefined Signals (1/2)

- 31 signals
- Every signal has a name.
 - begin with 'SIG'
 - SIGABRT: abort signal from abort()
 - SIGALRM: alarm signal from alarm()
- Actions of the default signal handler
 - terminate the process and generate a core (dump)
 - ignores and discards the signal (ignore)
 - suspends the process (suspend)
 - resume the process
- Default signal handler can be overridden.



Signal Sources



Predefined Signals (2/2)

```
esjung — esjung@hpclab:/usr/include — ssh hpclab — 84×32

#define SIGTERM      15
#define SIGSTKFLT    16
#define SIGCHLD      17
#define SIGCONT      18
#define SIGSTOP      19
#define SIGTSTP      20
#define SIGTTIN      21
#define SIGTTOU      22
#define SIGURG       23
#define SIGXCPU      24
#define SIGXFSZ      25
#define SIGVTALRM     26
#define SIGPROF      27
#define SIGWINCH     28
#define SIGIO        29
#define SIGPOLL       SIGIO
/*
#define SIGLOST      29
*/
#define SIGPWR       30
#define SIGSYS       31
#define SIGUNUSED    31

/* These should not be considered constants from userland. */
#define SIGRTMIN      32
#define SIGRTMAX      _NSIG

/*
 * SA_FLAGS values:
 *
 * SA_ONSTACK indicates that a registered stack_t will be used.
```

65,1 34%



Signal Generation

- Terminal-generated signals
 - CTRL-C → **SIGINT**
 - CTRL-Z → **SIGSTP** signal
- Hardware exceptions generate signals
 - divide by 0 → SIGFPE
 - invalid memory reference → SIGSEGV
- **kill()**
 - sends any signal to a process or process group
 - **need to be owner or super-user**
- Software conditions
 - SIGALRM: alarm clock expires
 - SIGPIPE: broken pipe
 - SIGURG: out-of-band network data



Handling of Signals

- Disposition or action:
Process has to tell the kernel “if and when this signal occurs, do the following.”
- Ignore the signal:
all signals can be ignored, **except SIGKILL and SIGSTOP**
- Let the default action apply:
most are to terminate process



Representative UNIX Signals (1/2)

- **SIGABRT**: generated by calling the **abort** function.
- **SIGALRM**: generated when a timer set with the **alarm** expires.
- **SIGCHLD**: whenever a process **terminates** or stops, the signal is sent to the parent.
- **SIGCONT**: this signal sent to a stopped process when it is **continued**.
- **SIGFPE**: signals an **arithmetic exception**, such as divide-by-0, floating point overflow, and so on
- **SIGILL**: indicates that the process has executed an **illegal hardware instruction**.
- **SIGINT**: generated by the terminal driver when we type the interrupt key and sent to all processes in the foreground process group.



Representative UNIX Signals (2/2)

- **SIGKILL**: can't be caught or ignored. a sure way to kill any process.
- **SIGPIPE**: if we write to a pipeline **but the reader has terminated**, SIGPIPE is generated.
- **SIGSEGV**: indicates that the process has made an invalid memory reference. (→ core dumped)
- **SIGTERM**: the termination signal sent by the kill(1) command by default.
- **SIGSTP**: **Cntl-Z** from the terminal driver which is sent to all processes in the foreground process group.
- **SIGUSR1**: user defined signal 1
- **SIGUSR2**: user defined signal 2



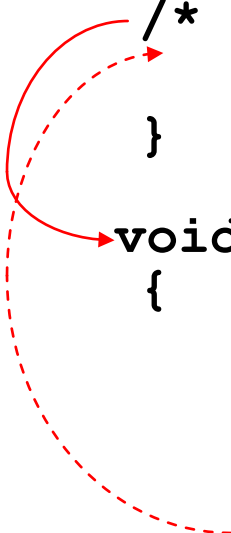
signal()

- Signal Handler Registration
- **signal(int signo, void(*func)())**
 - specify the action for a signal (*signo* → *func*)
 - *func*
 - SIG_IGN (ignore)
 - SIG_DFL (default)
 - user-defined function
 - Return: the previous *func*



Example

```
int main()  
{  
    signal( SIGINT, foo );  
    :  
    /* do usual things until SIGINT */  
    return 0;  
}  
void foo( int signo )  
{  
    :      /* deal with SIGINT signal */  
    return; /* return to program */  
}
```



A diagram consisting of two red curved arrows. The first arrow is solid and points from the `foo` function definition back to the `signal(SIGINT, foo);` line in the `main` function. The second arrow is dashed and points from the `return;` line in the `foo` function back to the `return 0;` line in the `main` function.



Example : alarm2.c (w/ handler) (1/2)

```
#include <stdio.h> // alarm2.c
#include <signal.h>

int alarmFlag=0;
void alarmHandler();

main( ) {
    signal(SIGALRM, alarmHandler);
    alarm(3);
    printf("Looping ...\n");
    while(!alarmFlag) {
        pause( );
    }
    printf("Loop ends due to alarm signal \n");
}

void alarmHandler( ) {
    printf("An alarm clock signal was received\n");
    alarmFlag = 1;
}
```



Example : alarm2.c (w/ handler) (2/2)

- Execution

```
[esjung@hpclab opensource]$ gcc -o alarm2 alarm2.c  
[esjung@hpclab opensource]$ ./alarm2  
Looping ...  
An alarm clock signal was received  
Loop ends due to alarm signal  
[esjung@hpclab opensource]$
```



SIGCHLD

- Whenever a process terminates or stops, the signal is sent to the [parent](#).
- When a child process is killed, it sends SIGCHLD signal to its parent process



Example: timelimit.c (1/3)

`$ timelimit N command` // perform “command” within N seconds

```
#include <stdio.h> // timelimit.c
#include <signal.h>

int delay;
void childHandler( );

main(int argc, char *argv[])
{
    int pid;

    sscanf(argv[1], "%d", &delay);
    signal(SIGCHLD, childHandler);

    pid = fork();
    if (pid == 0) { // child
        execvp(argv[2], &argv[2]);
        perror("Limit");
    } else { // parent
        sleep(delay);
        printf("Child %d exceeded limit and is being killed\n", pid);
        kill(pid, SIGINT);
    }
}
```



Example: timelimit.c (2/3)

```
childHandler( ) /* Executed if the child dies before the parent */
{
    int childPid, childStatus;
    childPid = wait(&childStatus);
    printf("Child %d terminated within %d seconds\n", childPid, delay);
    exit(0);
}
```



Example : timelimit.c (3/3)

- Execution

```
[esjung@hpclab opensource]$ ./timelimit 5 ls
7Process  alarm2.c  doatexit.c  lab3      pipe2.c  rdfile   timelimit  wtfile
alarm2    a.out      lab2      midterm  pipe.c   shfile.c  timilimit.c
Child 170561 terminated within 5 seconds
[esjung@hpclab opensource]$ ./timelimit 5 sleep 3
Child 170563 terminated within 5 seconds
[esjung@hpclab opensource]$ ./timelimit 5 sleep 100
Child 170573 exceeded limit and is being killed
[esjung@hpclab opensource]$
```



Multiple Signals

- If many signals of the *same* type are waiting to be handled (e.g. two `SIGINT`s), then most UNIXs will only deliver *one* of them.
 - the others are thrown away
- If many signals of *different* types are waiting to be handled (e.g. a `SIGINT`, `SIGSEGV`, `SIGUSR1`), they are not delivered in any fixed order.



The Reset Problem in early System V UNIX

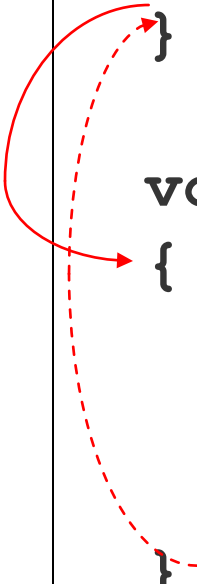
- In Linux (and many other UNIXs), the signal disposition in a process is **reset** to its **default action** **immediately after the signal has been delivered**.
- Must call `signal()` again to reinstall the signal handler function.



Reset Problem Example

```
int main()
{
    signal(SIGINT, foo);
    :
    /* do usual things until SIGINT */
}

void foo(int signo)
{
    signal(SIGINT, foo); /* reinstall */
    :
    return;
}
```

A red dashed line with arrows at both ends forms a loop. One arrow points from the closing brace of the `foo` function back to the `signal(SIGINT, foo);` line inside the `foo` function. The other arrow points from the closing brace of the `main` function back to the `signal(SIGINT, foo);` line in the `main` function, illustrating the recursive nature of the signal handler.

Reset Problem

```
void ouch( int sig )
{
    printf( "OUCH! - I got signal %d\n", sig );
    (void) signal(SIGINT, ouch);
}

int main()
{
    (void) signal( SIGINT, ouch );
    while(1)
    {
        printf("Hello World!\n");
        sleep(1);
    }
}
```

To keep catching the signal with this function, must call the *signal* system call again.

Problem: from the time that the interrupt Function starts to just before the signal handler is re-established the signal will not be handled.

If another SIGINT signal is received during this time, default behavior will be done, i.e., program will terminate.



Re-installation may be too slow!

- There is a (very) small time period in `foo()` when a new `SIGINT` signal will cause the default action to be carried out -- process termination.
- **POSIX, BSD** signal functions **solve** it (and some other later UNIXs)



Modification in BSD 4.x signal environment

- Persistent Handlers
 - Signal handlers remain installed even after the signal occurs and do not need to be explicitly reinstalled.
- Masking
 - Signals are blocked for the duration of a signal handler (*i.e.* recursive signals are not normally allowed).
 - A "signal mask" can be set to block most signals during critical regions.
- Signal handlers normally remain installed during and after signal delivery.



kill(), raise()

```
#include <sys/types.h>
#include <signal.h>

int kill(pid_t pid, int signo);
int raise(int signo);
```

Both return: 0 if OK, 1 on error

- **kill** - sends a signal to a process or a group of process
- **raise** - function allows a process to send a signal to itself



kill()

- pid means
 - $\text{pid} > 0$: signal to the process whose process ID is pid
 - $\text{pid} == 0$: signal to the processes whose process group ID equals that of sender
 - $\text{pid} < 0$: signal to the processes whose process group ID equals *abs.* of pid
 - $\text{pid} == -1$: unspecified (used as a broadcast signal in SVR4, 4.3 + BSD)
- Permission to send signals
 - The super-user can send a signal to any process.
 - The real or effective user ID of the sender has to equal the real or effective user ID of the receiver.



alarm()

```
#include <unistd.h>
```

```
unsigned int alarm (unsigned int seconds);
```

```
Returns: 0 or number of seconds until previously set alarm
```

- alarm() **sets a timer** to expire at a specified time in future.
 - when timer expires, SIGALRM signal is generated,
 - default action of the signal is to terminate the process.
- Only one alarm clock per process
 - previously registered alarm clock is replaced by the new value.
- if alarm(0), a previous unexpired alarm is cancelled.



pause()

```
#include <unistd.h>
```

```
int pause (void);
```

Returns: -1 with errno set to EINTR

- suspends the calling process until a signal is caught.
- returns only if a signal handler is executed and that handler returns.
 - If signal handler is not registered, just quit
 - If signal handler is registered, return after the handler is processed.



abort()

```
#include <stdlib.h>

void abort(void);
    This function never returns
```

- Causes abnormal program termination.
- This function sends the SIGABRT signal to the process.
- SIGABRT signal handler to perform any cleanup that it wants to do, before the process terminated.



sleep()

```
#include <signal.h>
```

```
unsigned int sleep(unsigned int seconds) ;
```

```
    Returns: 0 or number of unslept seconds
```

- This function causes the calling process to be suspended until either
 - The amount of wall clock time specified by second has elapsed (returns 0)
 - A signal is caught by the process and the signal handler returns (returns the number of unslept seconds)



Process group

- Group of related processes that should receive a common signal for certain events
- Identified by Process group ID
- Example
 - `grp = setpgp();`
 - Initialize process group id, the same as its pid
 - After `fork()`, child process will have parent's process group id



setpgrp()

```
#include <signal.h>
#include <stdio.h>

main() {
    register int i;
    setpgrp();
    for(i=0;i<10;i++){
        if(fork()==0){
            /* child process */
            if(i&1)
                setpgrp();
            printf("pid = %d pgrp= %d\n", getpid(),getpgrp());
            pause();
        }
    }
    kill(0,SIGINT);
}
```



Results

```
[esjung@hpclab opensource]$ ./a.out  
pid = 100672 pgrp= 100671  
pid = 100673 pgrp= 100673  
pid = 100674 pgrp= 100671  
pid = 100676 pgrp= 100671  
pid = 100675 pgrp= 100675  
pid = 100677 pgrp= 100677  
pid = 100678 pgrp= 100671  
pid = 100679 pgrp= 100679  
pid = 100680 pgrp= 100671
```



Treatment of Signals

- Sending a signal to a process
 - `issig()` : check for receipt of signals
- Handling a signal
 - `psig()` : handle signals after recognizing their existence



How to Send a Signal to Process

- Send a Signal

- sets a bit in the signal field in process table entry.
- process can remember different type of signal.
- process can not remember how many signals it receives of particular type.
- u area contains an array of signal-handler fields.
kernel stores the address of the user-function in the field.

- Check for Signal **when**

- about to return from kernel mode to user mode
- enters or leaves the sleep state at a suitably low scheduling priority.

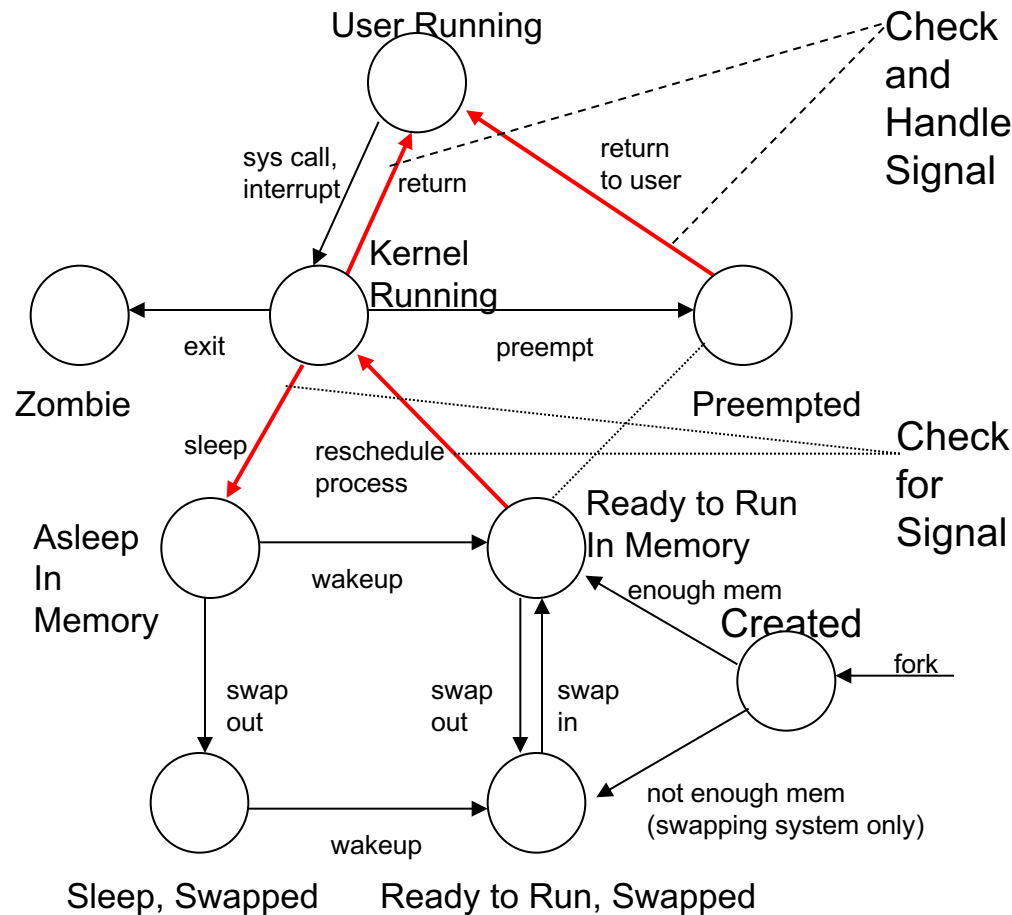
- Handle Signals

- only when returns from kernel mode to user mode



How to Send a Signal to Process(Cont.)

- no effect on a process running in the kernel mode
- a process never executes in user mode before handling outstanding signals



Algorithm for issig

```
algorithm issig          /* test for receipt of signals */
input : none
output : true, if process receives signals that it does not ignore
        false otherwise
{
    while(received signal field in process table entry not 0)
    {
        find a signal number sent to the process;
        if ( signal is death of child )
        {
            if (ignoring death of child signals)
                free process table entries of zombie children;
            else if (catching death of child signals)
                return(true);
        }
        else if (not ignoring signal)
            return(true);
        turn off signal bit in received signal field in process table;
    }
    return(false);
}
```



Algorithm for Handling Signal

1. determines **signal type**
2. **turns off signal bit** in the process table entry
3. if receives a signal to **ignore**
 - continues as if the signal has never occurred.
4. If signal handling function is set to its **default** value,
 - kernel dumps core image for signals that imply something is wrong with process and exit
 - kernel does not dump core for signals that do not imply error.
5. If receives **a signal to catch**,
 - accesses the user saved register context, and find the program counter and stack pointer
 - clears the signal handler field in u area(*undesirable side-effects*)
 - kernel creates a new stack frame and writes a program counter and stack pointer from user saved register context
 - kernel changes the user register context: program counter to the address of signal catcher function, and stack pointer to account for the growth of the user stack.



Algorithm for psig

```
algorithm psig          /* handle signals after recognizing their
    existence */
input : none
output : none
{
    get signal number set in process table entry;
    clear signal number in process table entry;
    if (user had called signal sys call to ignore this signal)
        return;          /*done*/
    if (user specified function to handle this signal)
    {
        get user virtual address of signal catcher stored in u area;
        /*the next statement has undesirable side-effects*/
        clear u area entry that stored address of signal catcher;
        modify user level context;
            artificially create user stack frame to mimic
            call to signal catcher function;
        modify system level context;
            write address of signal catcher into program
            counter filed of user saved register context;
        return;
    }
}
```



Algorithm for psig (cont.)

```
if (signal is type that system should dump core image of
process)
{
    create file named "core" in current directory;
    write contents of user level context to file "core";
}
invoke exit algorithm immediately;
}
```

